



https://loma-aimotion.github.io/ml_visualizations/MiB_schiendorfer.pdf

Siehst du den Fehler? Wie funktionieren eigentlich Sichtprüfungen mit KI in der Produktion? Prof. Dr. Alexander Schiendorfer



KI-basierte Optimierung in der
Automobilproduktion



Technische Hochschule
Ingolstadt

Forschungsprofessur Schiendorfer

Doktoranden und PostDocs

- Lukas Lodes, MSc
- Kristina Dachtler, MEng
- Dr. Karima Outafraout
- Pauline Steffel, MSc
- Gheorghe Lisca, MSc
- Mert Alagözlü, MEng

- F+E Partner
 - Valeo
 - plus10 GmbH
 - Convergent-IT
 - PSI Software
 - Büchl Entsorgungswirtschaft



Key facts „Almotion Bavaria“ aimotion.de

Bayerischer KI-Mobilitätsknoten an der THI



- 20 Forschungsprofessuren
 - *Schwerpunkte:* Autonome Mobilität (Fahren, Fliegen), **KI-gestützte (Automobil)produktion**
 - *Zielsetzung:* Aufbau von bis zu 120 Forschende durch Drittmittelforschung (dzt. 75)
- Im Gebäude K (Neuer Dalwigk)
 - 4.000 m² Fläche, Seminarräume, KI-Labore
 - Ausbau Informatik, KI im Mobilitätsknoten



Institut Almotion Bavaria in Zahlen (Stand: 31.12.2023)



Institute des KI-Mobilitätsknotens:



KI-Ökosystem:

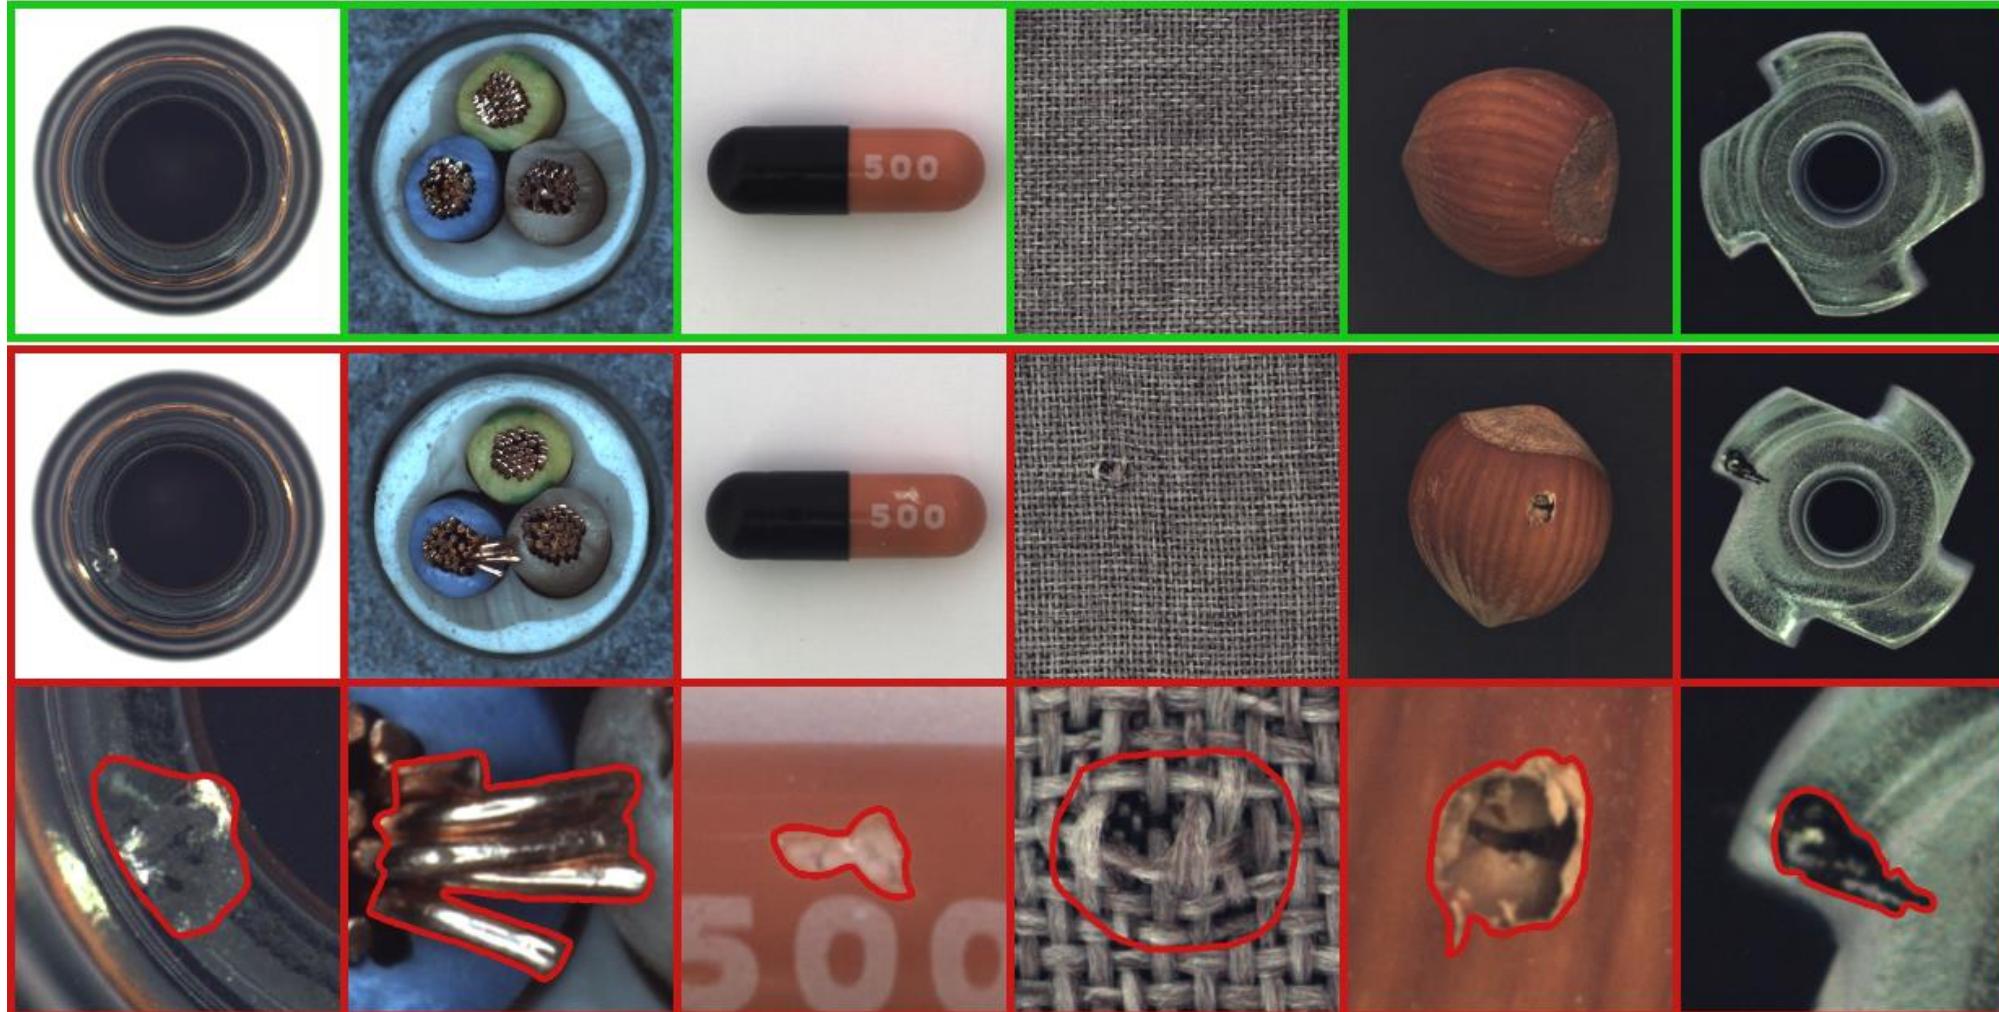


AI in Automotive Production



KI-basierte Bildverarbeitung

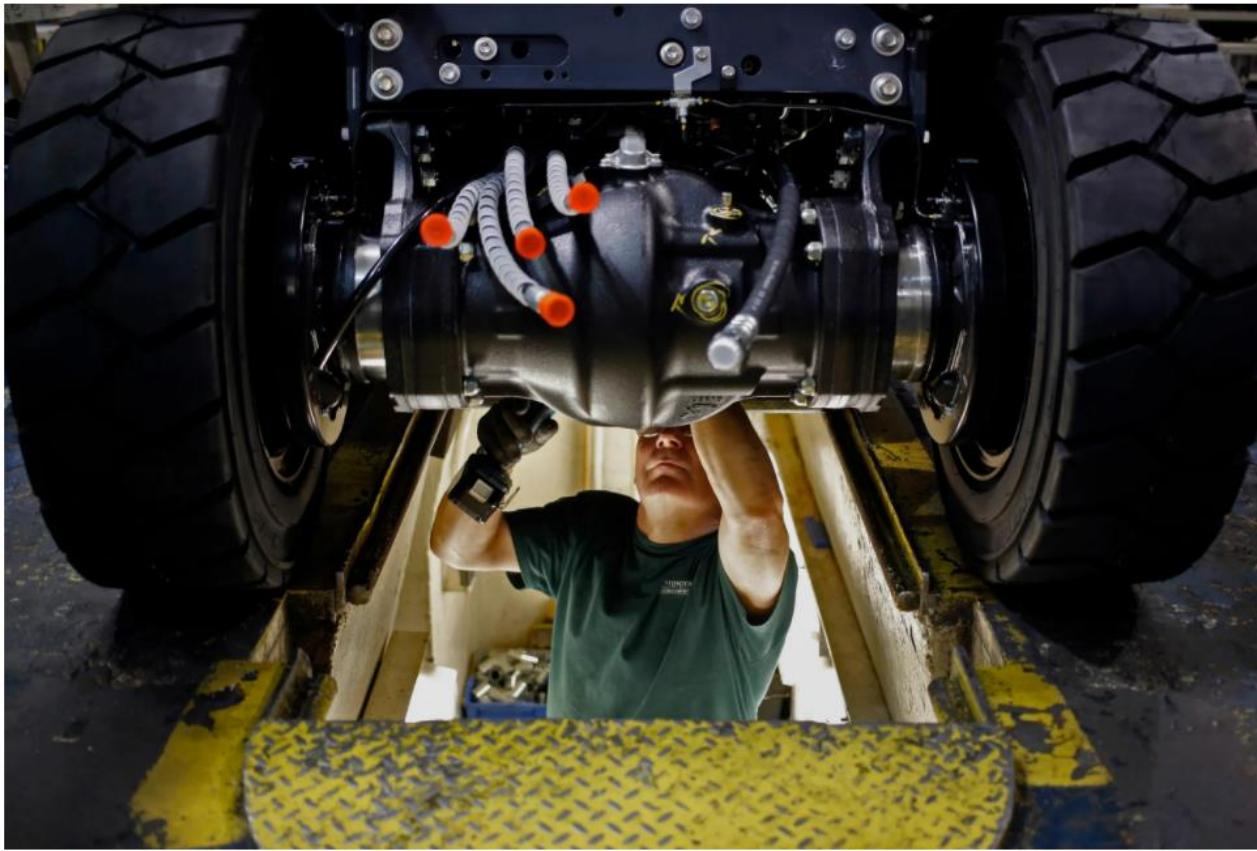
Visual anomaly detection in industrial use cases



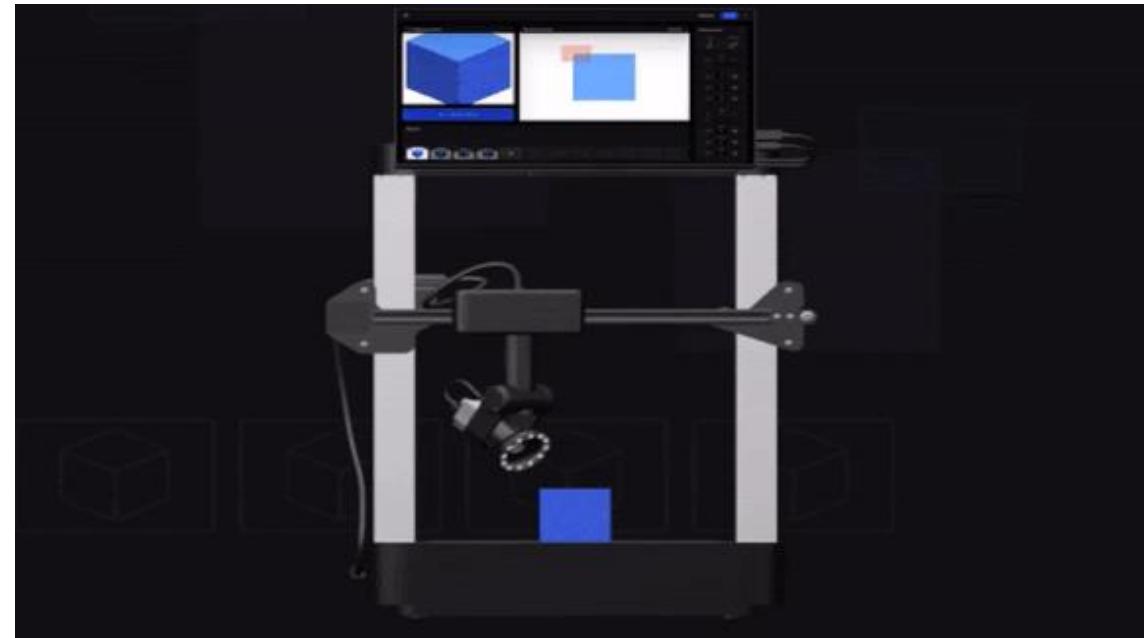
<https://www.mvtec.com/company/research/datasets/mvtec-ad>

In These Factories, Inspector Robot Will Check Your Work

Artificially intelligent camera systems look for defects and misplaced parts in many industries. The coronavirus pandemic makes them extra useful.



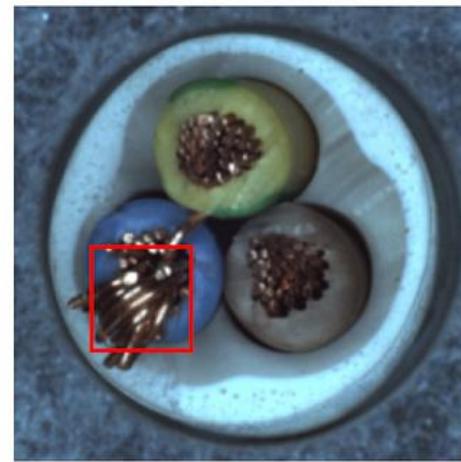
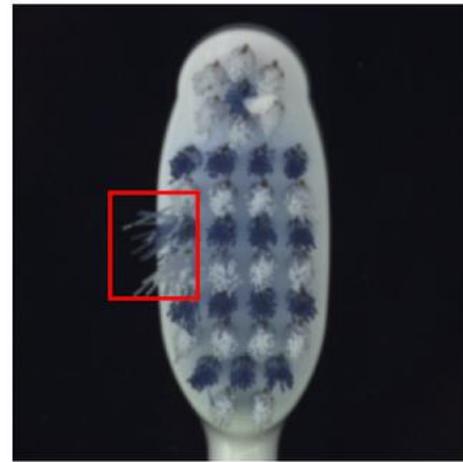
A Toyota assembly plant in Indiana is preparing to deploy a robotic system that moves a camera around an object, to ensure employees are using the correct parts. PHOTOGRAPH:
LUKE SHARRETT/BLOOMBERG/GETTY IMAGES



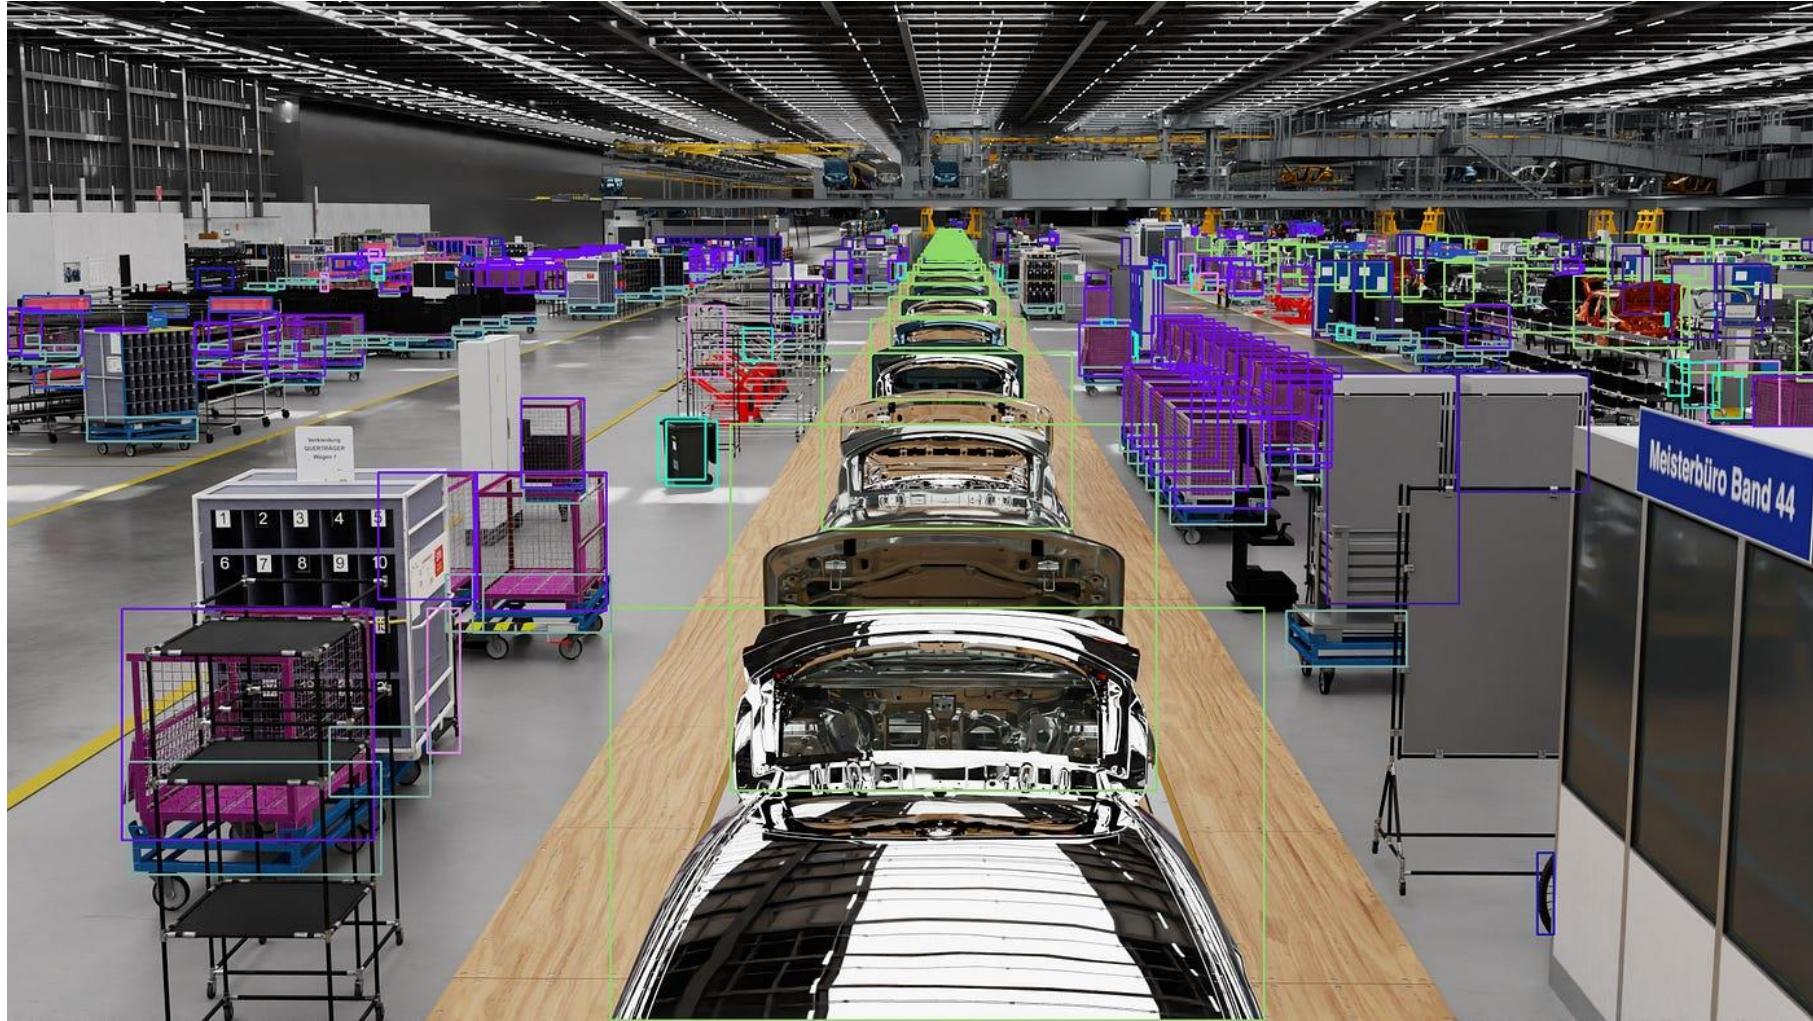
INSTRUMENTAL

<https://www.wired.com/story/factories-inspector-robot-check-work/>

Defect Detection

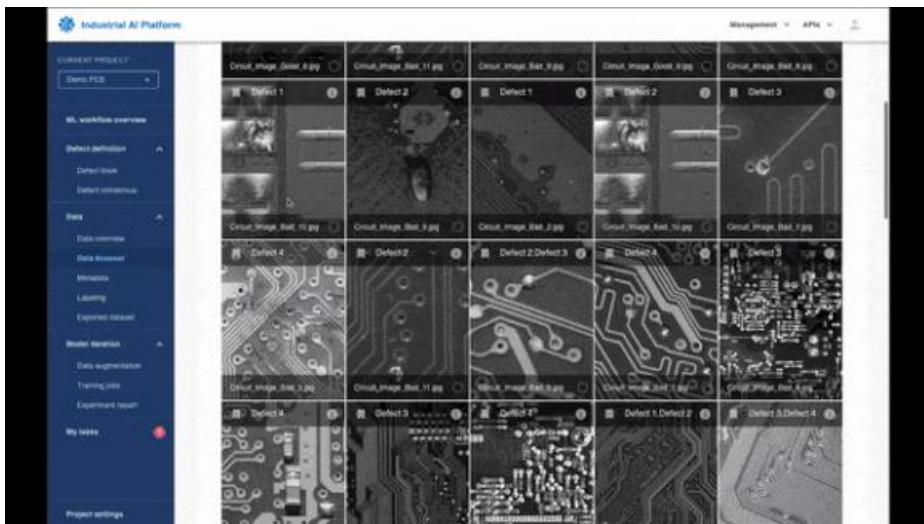


Tracking the state of a jobshop in production

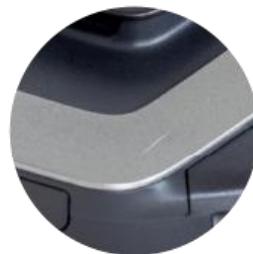


<https://sordi.ai/>

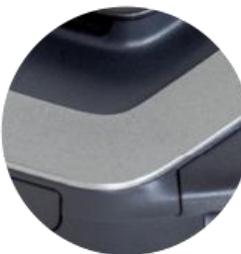
You'll work with AI systems!



Unambiguous defect ($y=1$)
All labelers will agree



Ambiguous whether to classify as defect
Even expert labelers disagree



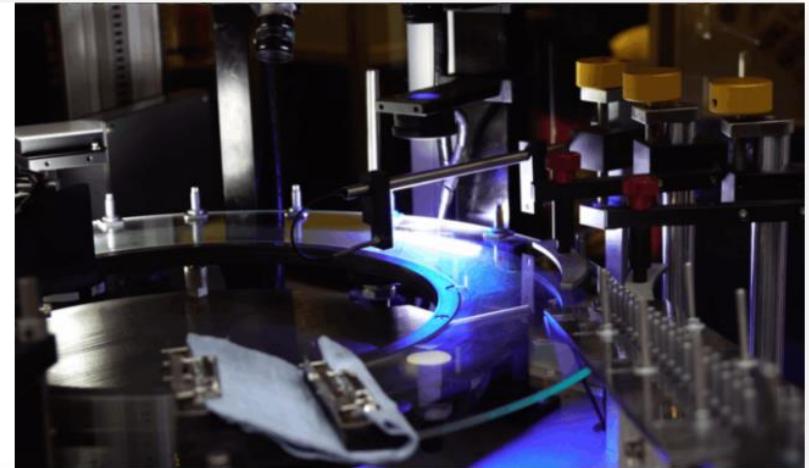
Unambiguous no defect ($y=0$)
All labelers will agree



Manufacturing

Every scratch, chip, or defective part matters to the bottom line. Identifying defects correctly and consistently is critical to production success.

Inspectors are constantly challenged with defining defects randomly. Landing AI provides an integrated, end-to-end visual inspection platform that effectively manages your data and builds solutions you can rely on.



Rather than holding the training set fixed and trying to improve the model, we hold the model fixed and help manufacturers improve the training set. We've found that this approach leads to faster progress in production settings.

<https://landing.ai/industries/#automotive>

Other application: Visual quality inspections at Audi

10/15/18 | Ingolstadt | Company

Audi optimizes quality inspections in the press shop with artificial intelligence



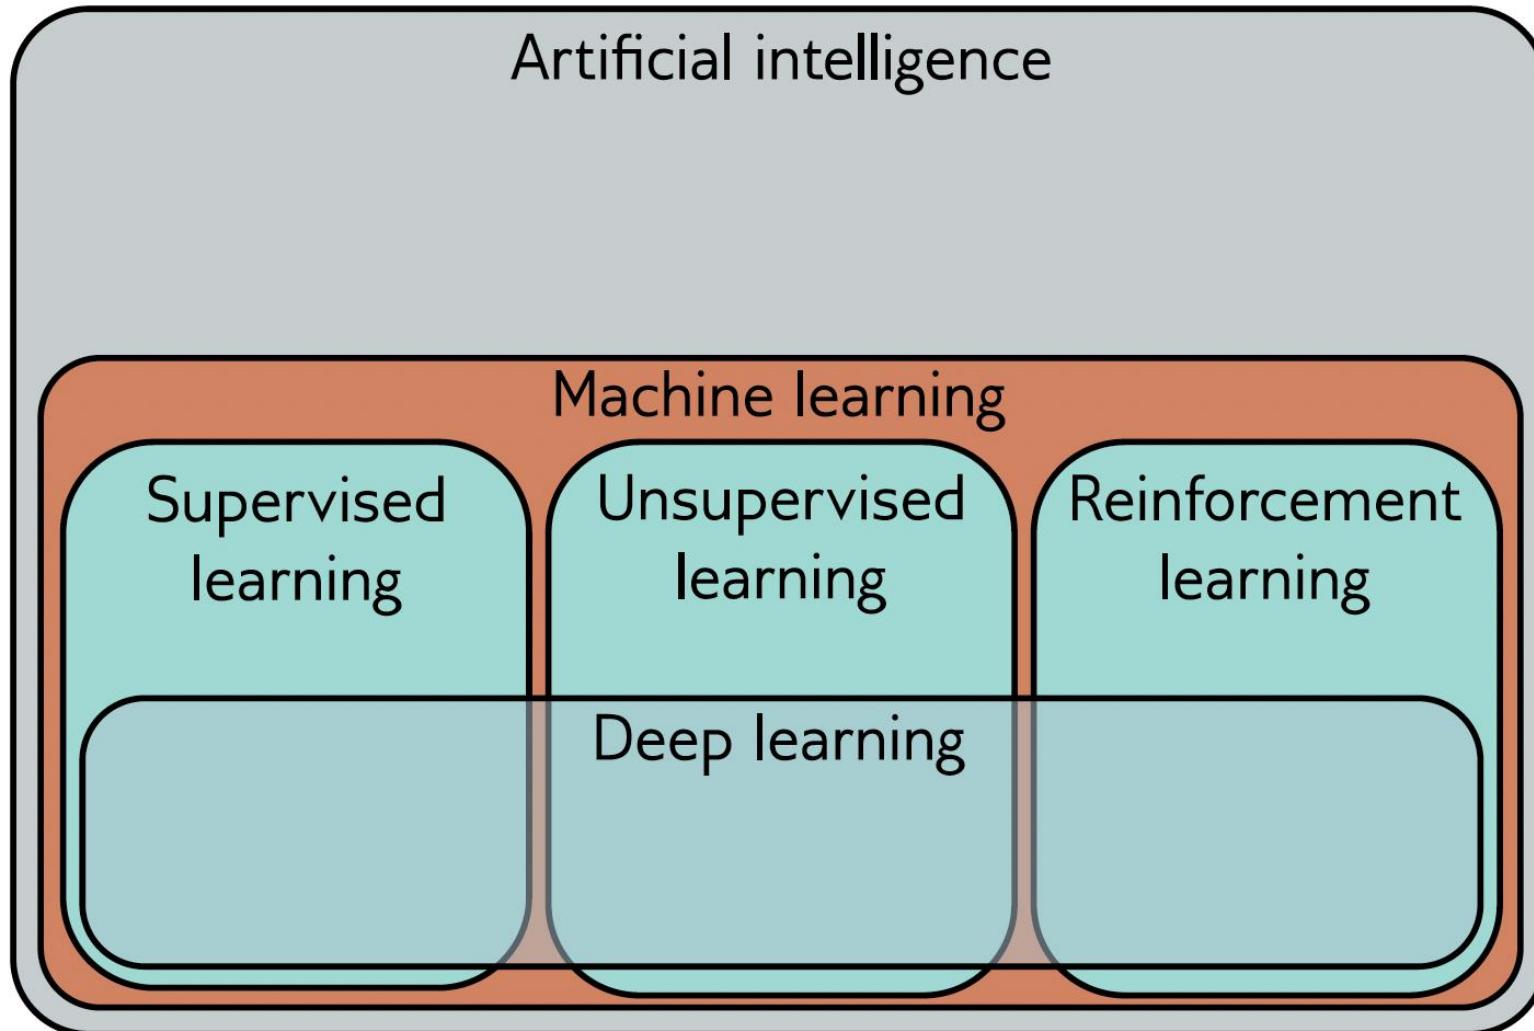
- First application case goes into series development
- Crack detection in the press shop to be automated with machine learning



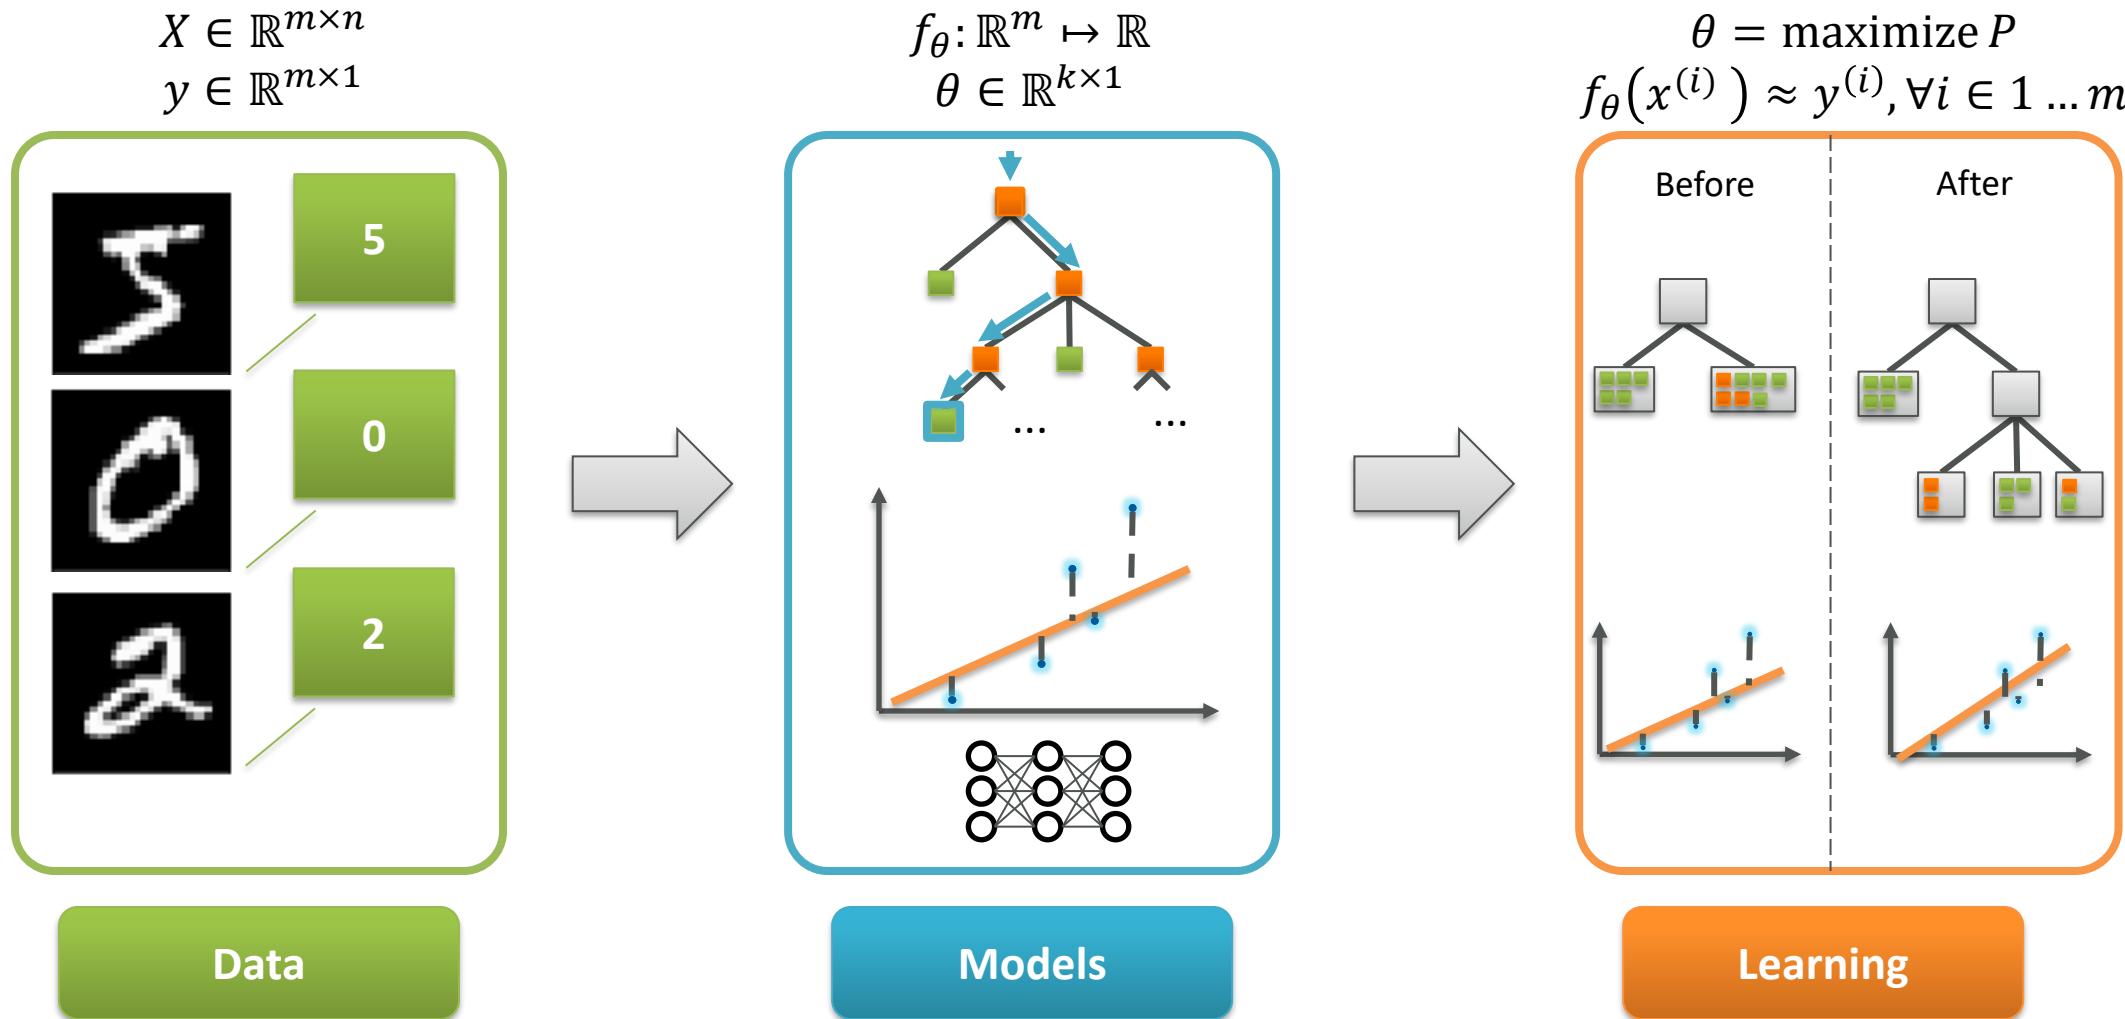
<https://www.audi-mediacenter.com/en/press-releases/audi-optimizes-quality-inspections-in-the-press-shop-with-artificial-intelligence-10847>

What is Deep Learning?

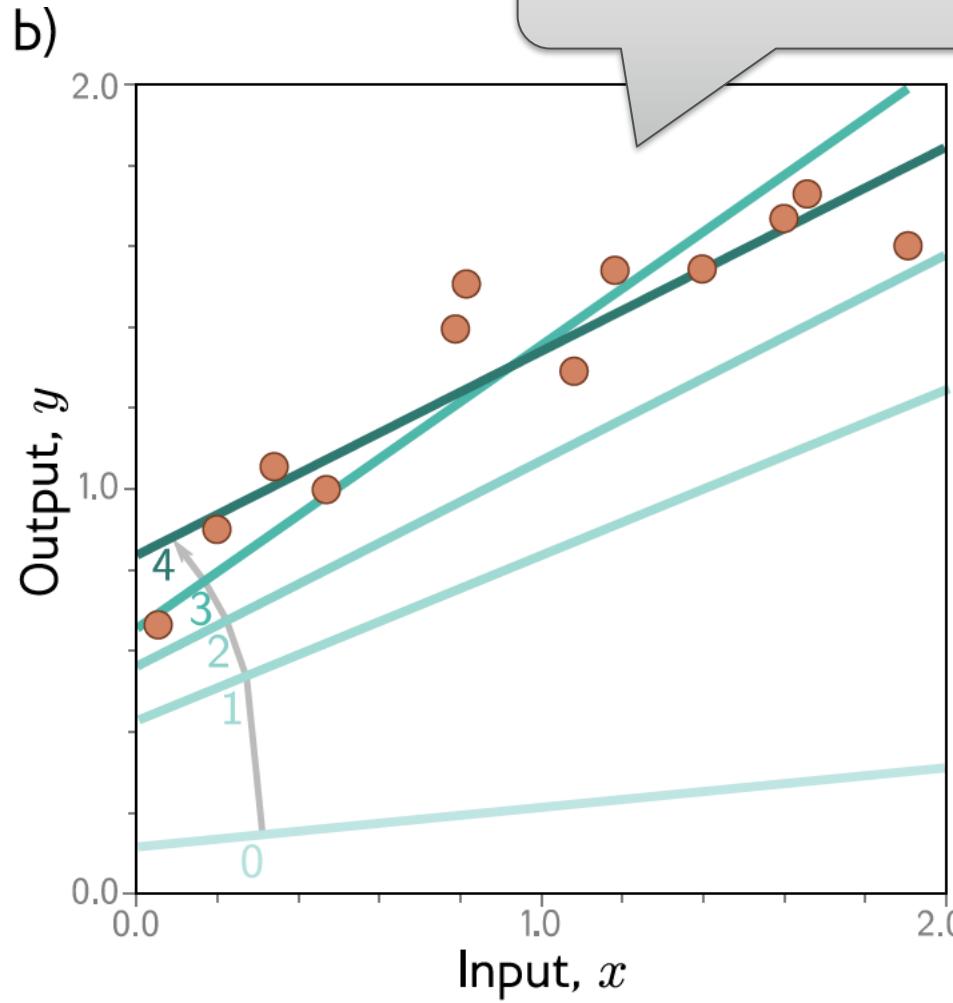
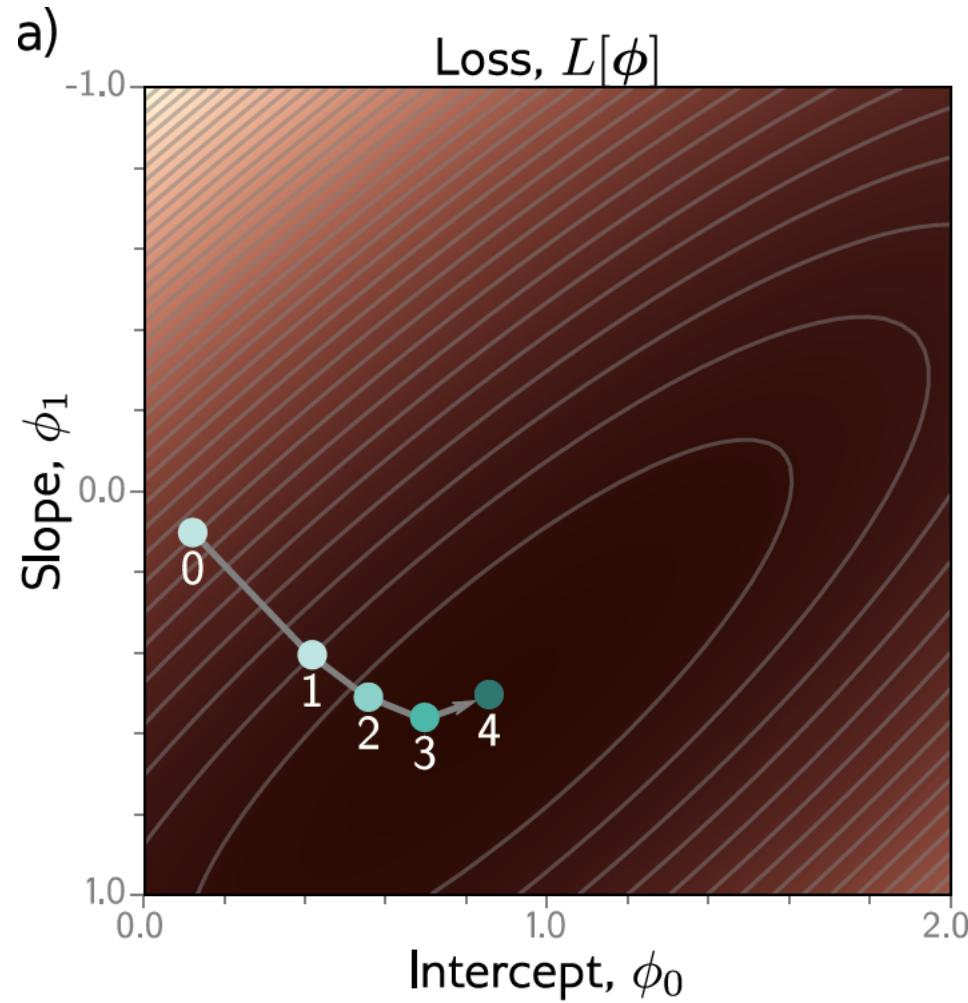
Categorization of AI



The three pillars of all machine learning systems



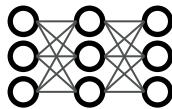
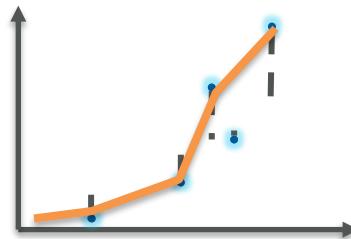
Example of Non-Deep ML: Linear Regression



In deep learning:

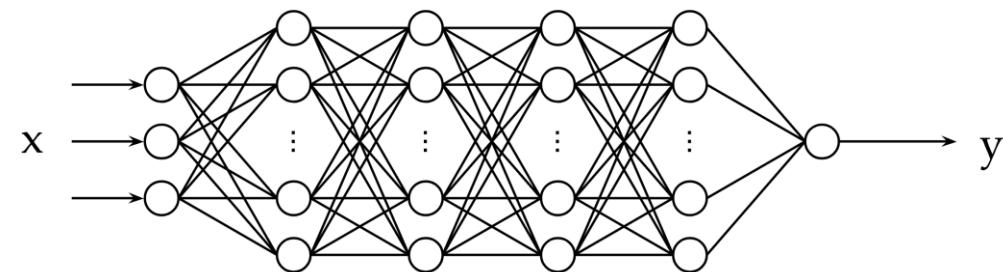
$$f_{\theta}: \mathbb{R}^m \mapsto \mathbb{R}$$

$$\theta \in \mathbb{R}^{k \times 1}$$



Models

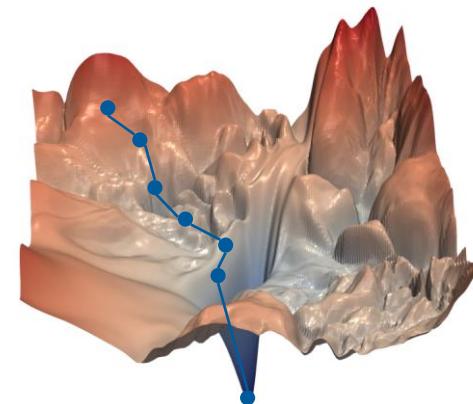
Models are "neural networks" with many parameters, stacked in many layers -> **universal function approximators**



Learning is some form of **gradient descent**, where a loss function between current and desired target values is minimized.

$$\theta = \text{maximize } P$$

$$f_{\theta}(x^{(i)}) \approx y^{(i)}, \forall i \in 1 \dots m$$

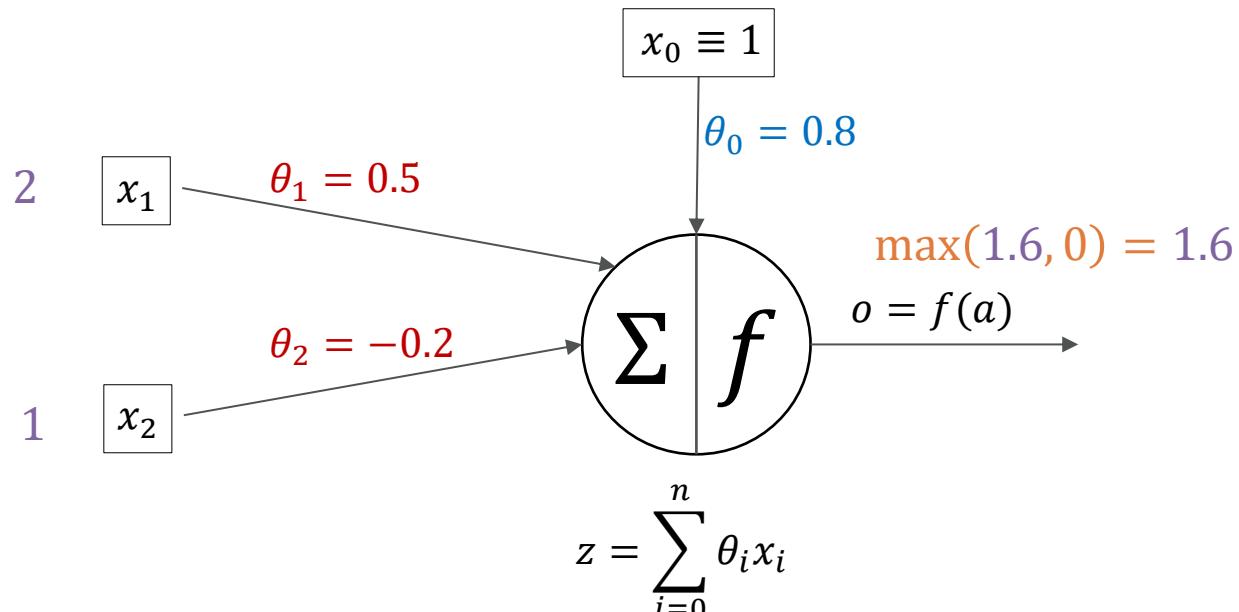


Learning

<https://arxiv.org/pdf/1712.09913.pdf>

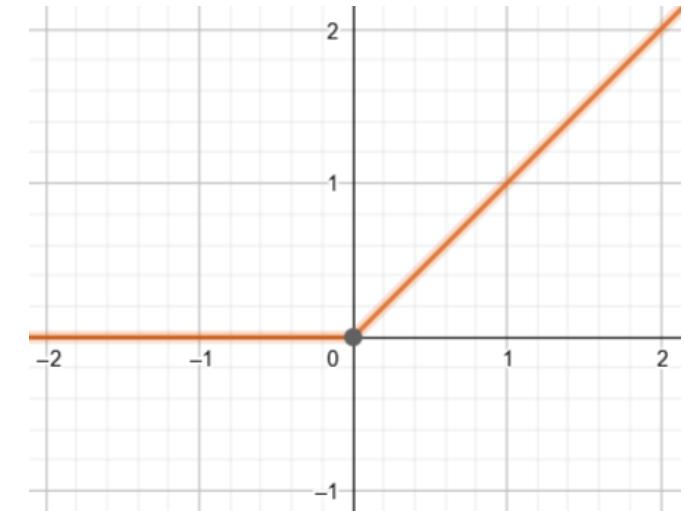
Inside a Neural Network

Example (single) Artificial Neuron



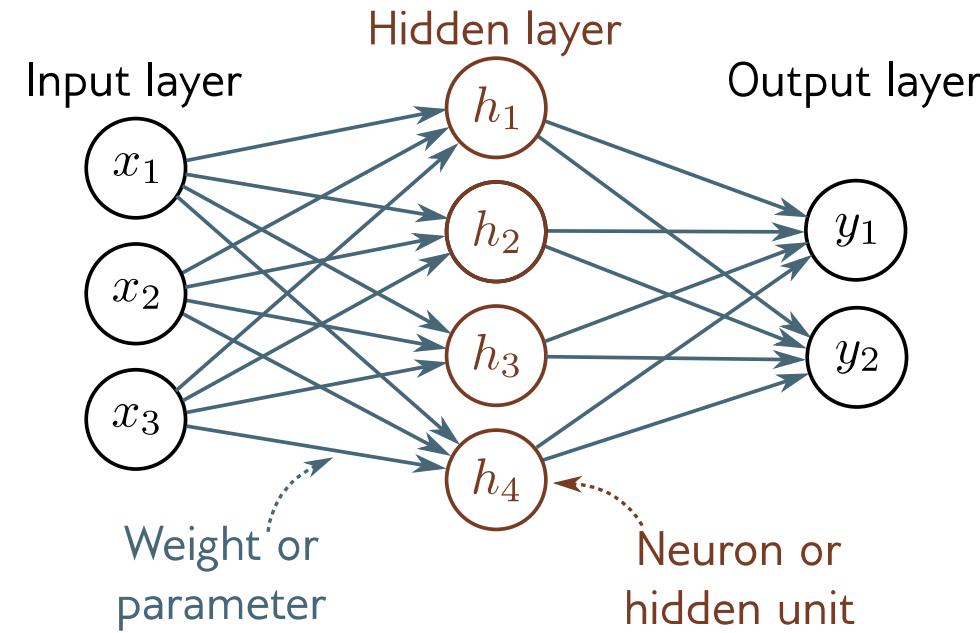
$$2 \cdot 0.5 + 1 \cdot (-0.2) + 0.8 = 1.6$$

x_i :	inputs	θ_i :	weights
θ_0 :	bias weight	$a(z)$:	activation function
Z :	pre-activation	o :	output



$$f(z) = \max(z, 0)$$

Nomenclature



- Y-offsets = **biases**
- Slopes = **weights**
- Everything in one layer connected to everything in the next = **fully connected network**
- No loops = **feedforward network**
- Values after ReLU (activation functions) = **activations**
- Values before ReLU = **pre-activations**
- One hidden layer = **shallow neural network**
- More than one hidden layer = **deep neural network**
- Number of hidden units \approx **capacity**

https://loma-aimotion.github.io/ml_visualizations/shallow_neural_network.html

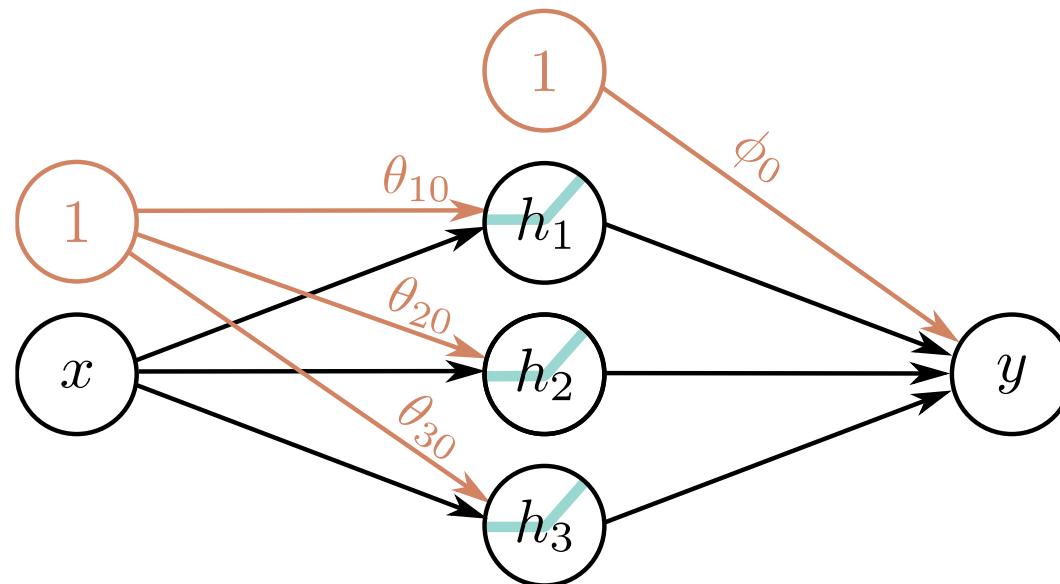
Depicting neural networks

$$h_1 = a[\theta_{10} + \theta_{11}x]$$

$$h_2 = a[\theta_{20} + \theta_{21}x]$$

$$h_3 = a[\theta_{30} + \theta_{31}x]$$

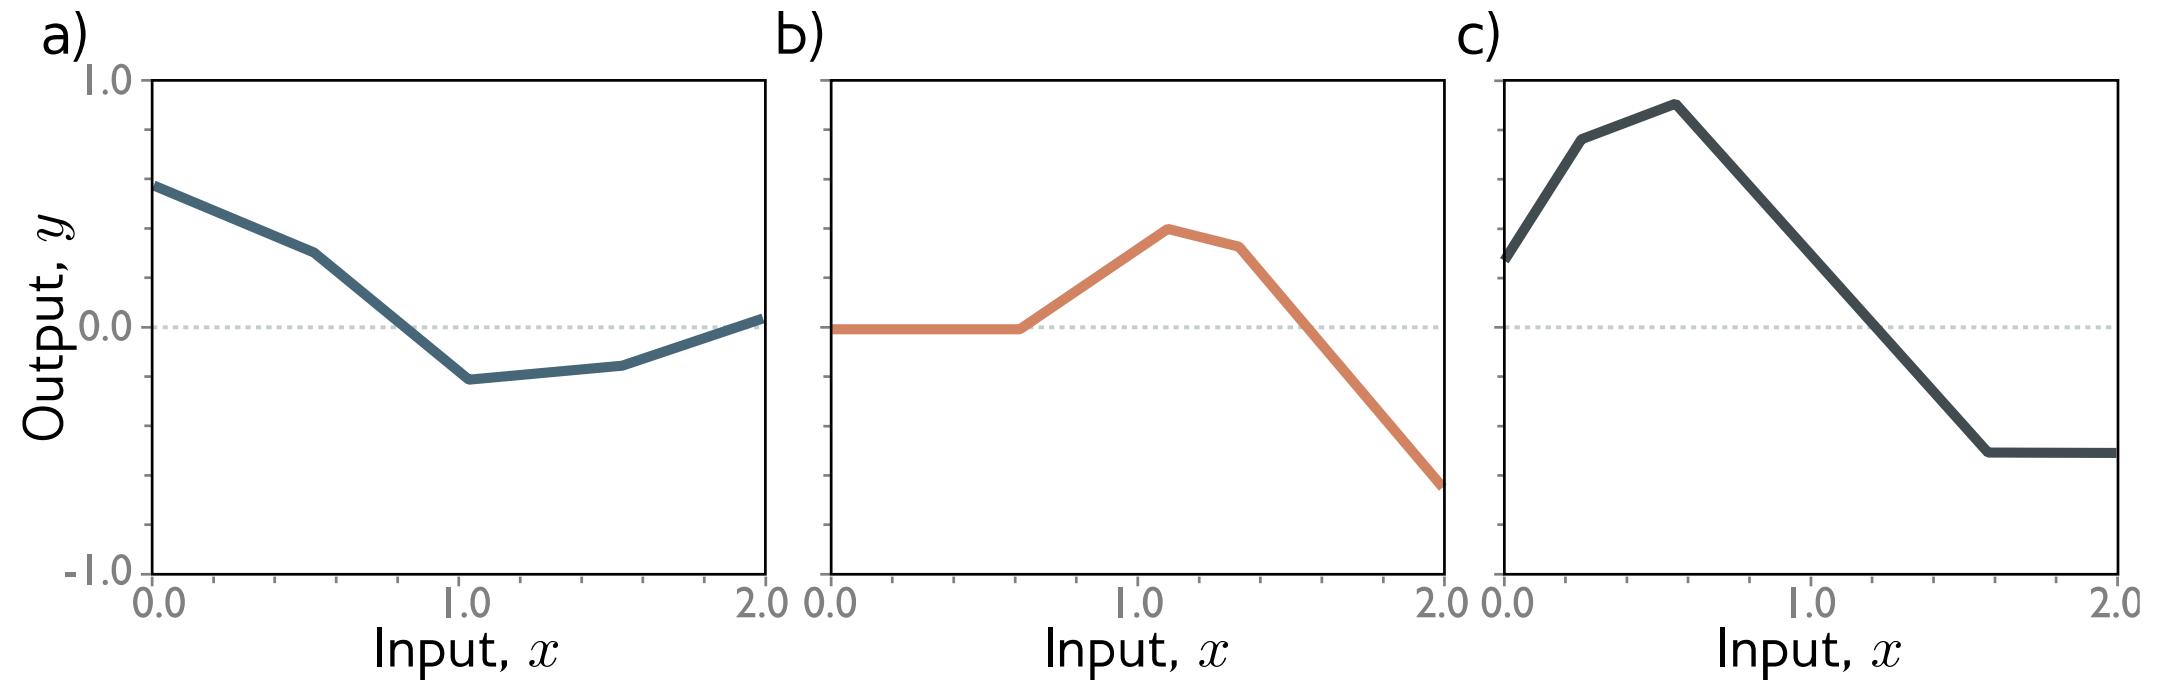
$$y = \phi_0 + \phi_1 h_1 + \phi_2 h_2 + \phi_3 h_3$$



Each parameter multiplies its source and adds to its target

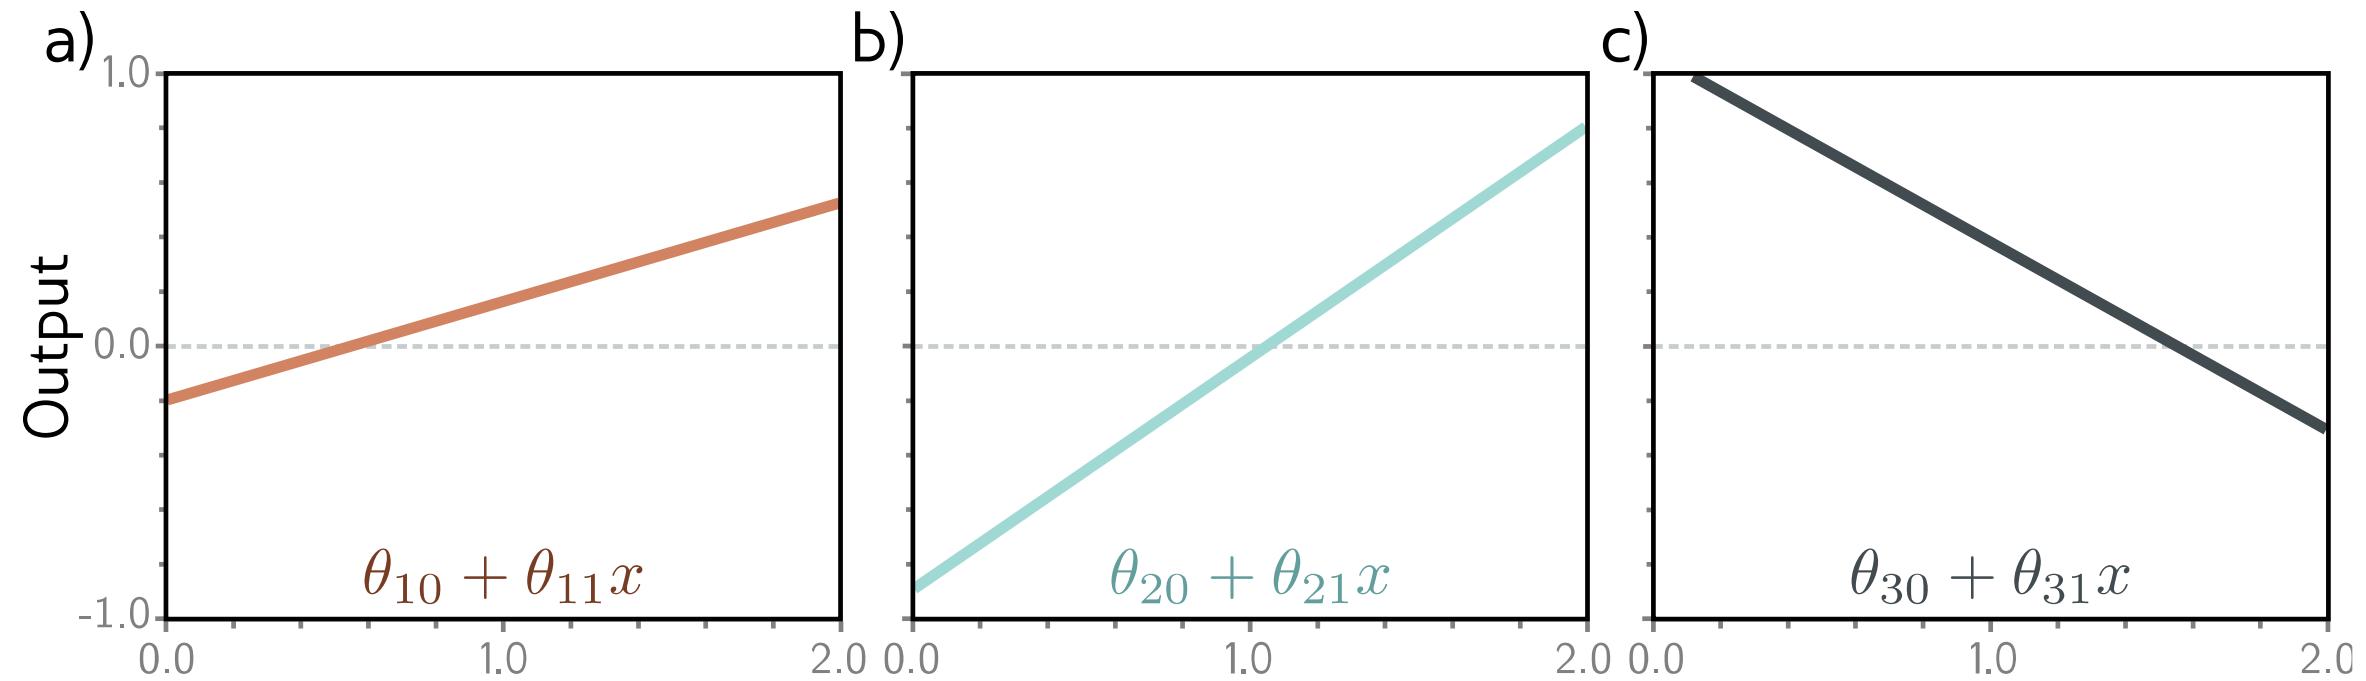
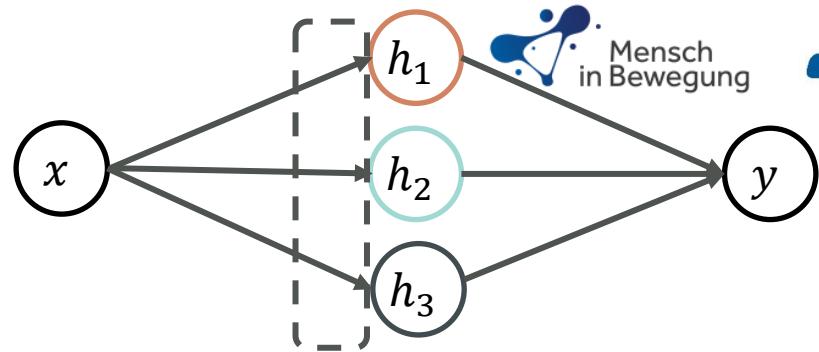
Example shallow network

$$y = \phi_0 + \phi_1 a[\theta_{10} + \theta_{11}x] + \phi_2 a[\theta_{20} + \theta_{21}x] + \phi_3 a[\theta_{30} + \theta_{31}x].$$



Piecewise linear functions with three joints

1. compute three linear functions

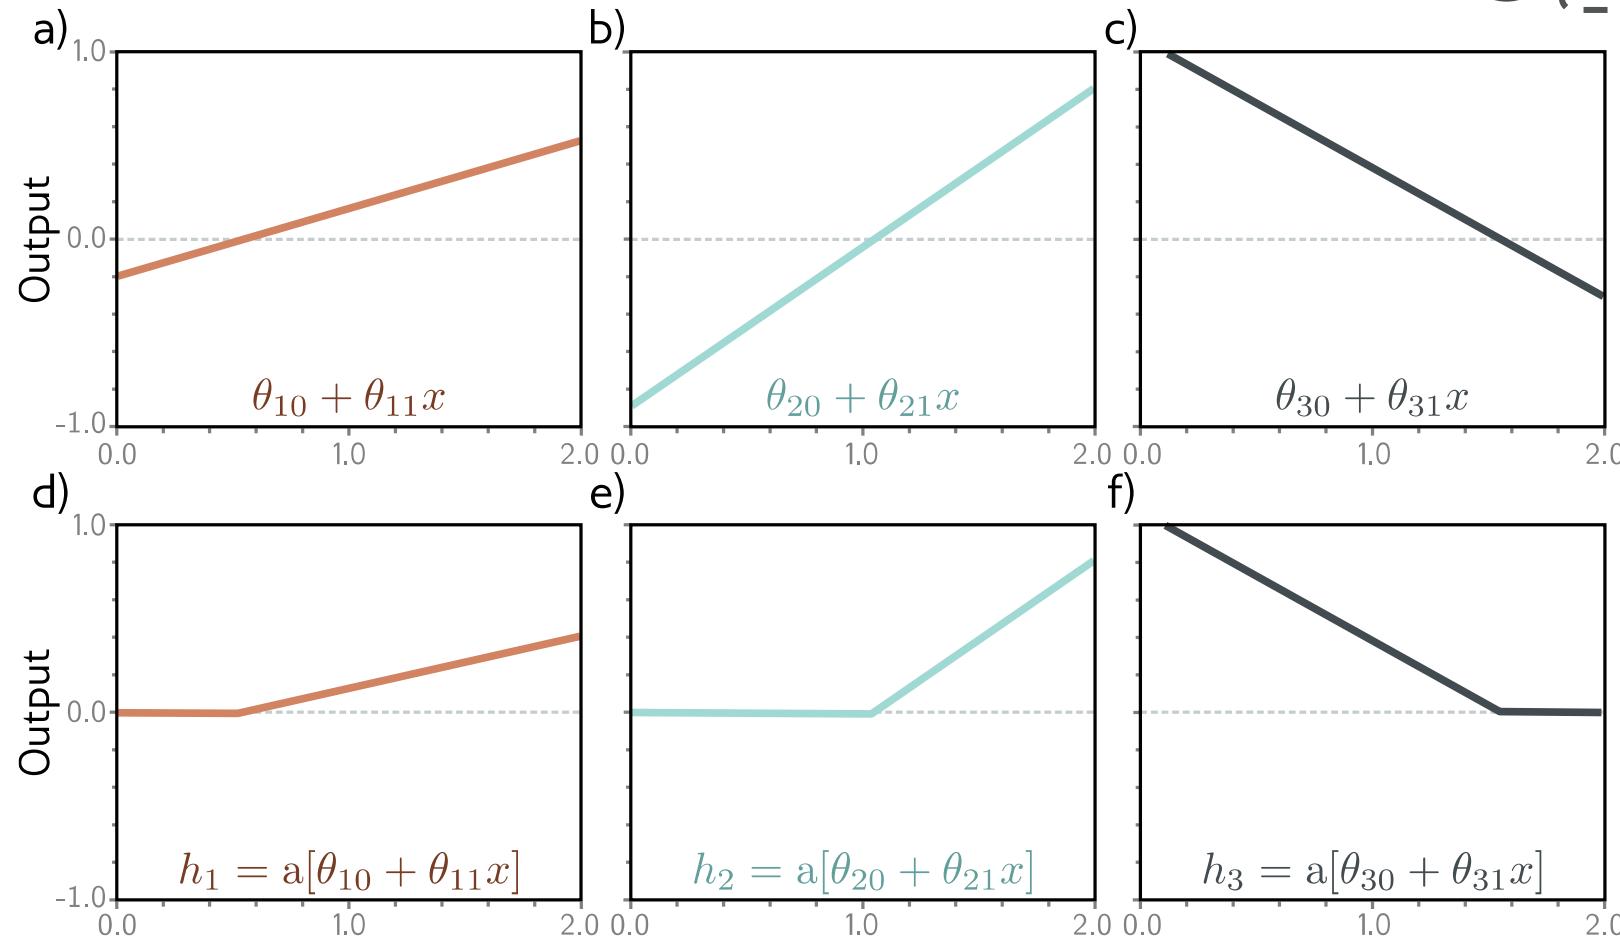
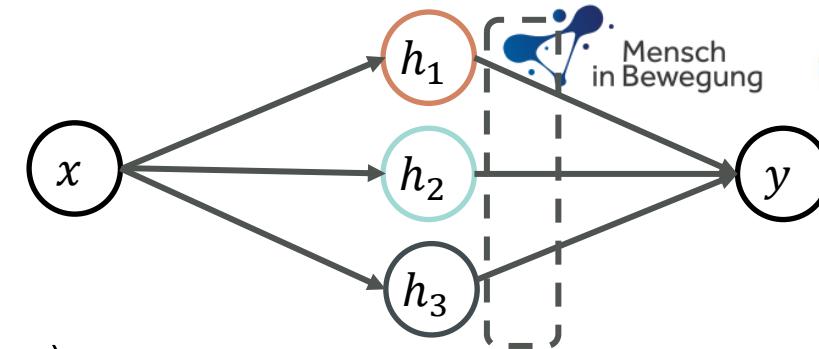


2. Pass through ReLU functions

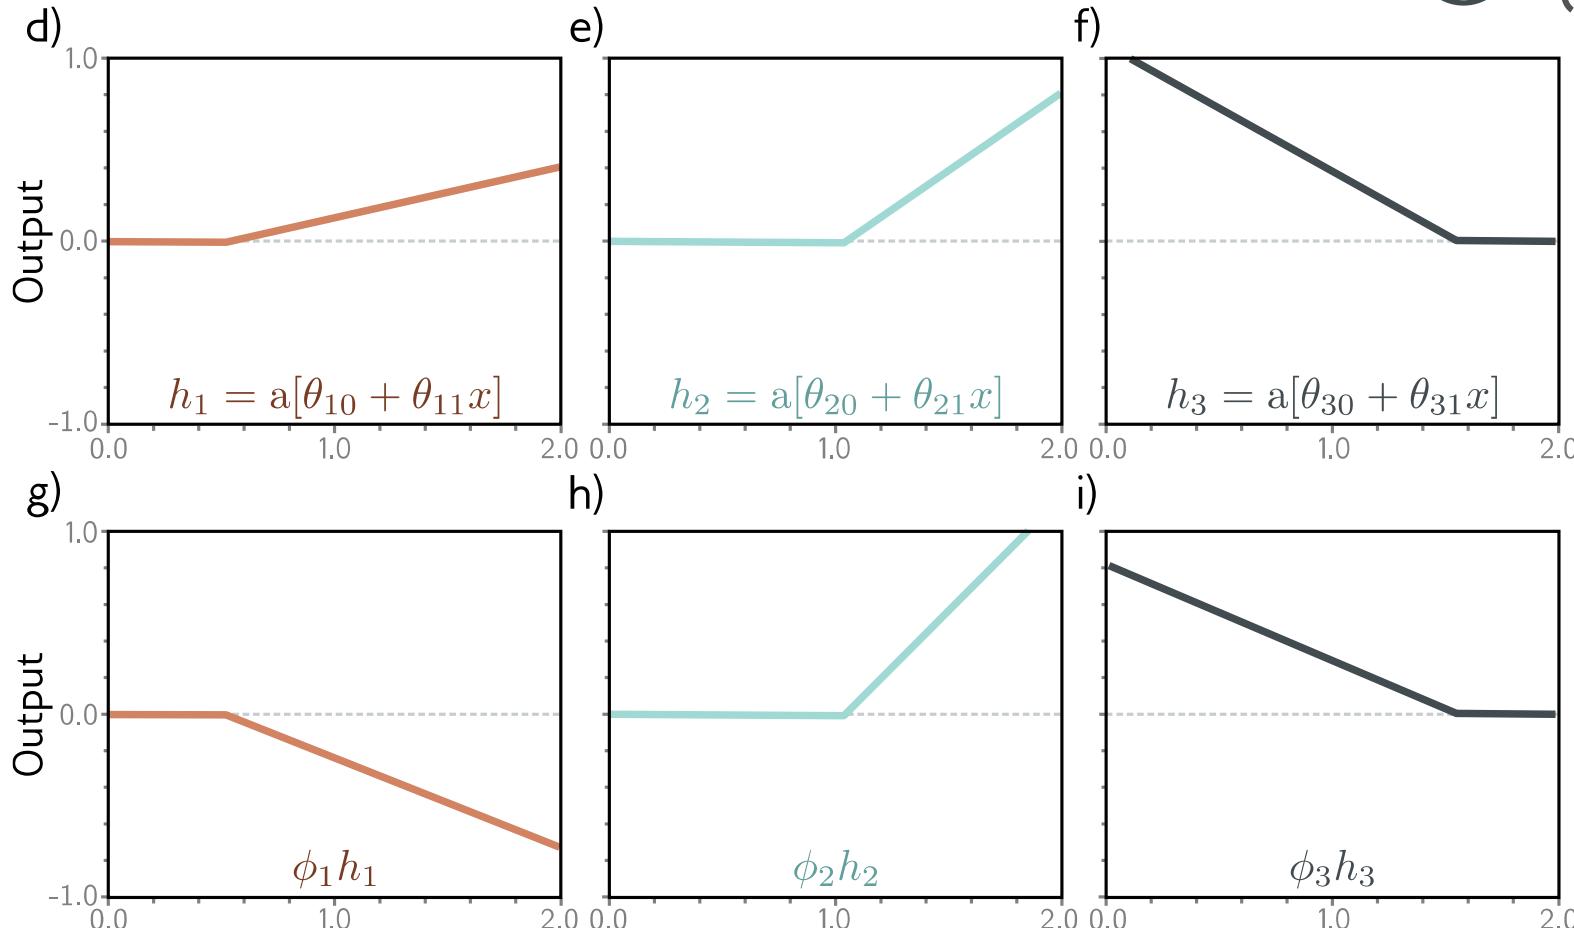
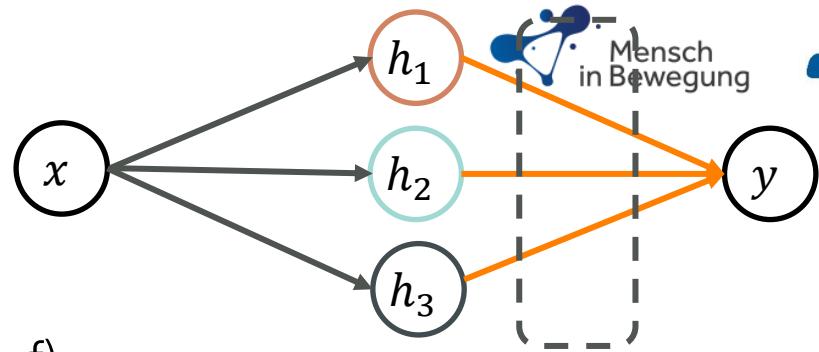
$$h_1 = a[\theta_{10} + \theta_{11}x]$$

$$h_2 = a[\theta_{20} + \theta_{21}x]$$

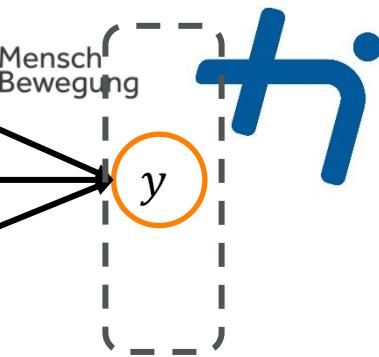
$$h_3 = a[\theta_{30} + \theta_{31}x],$$



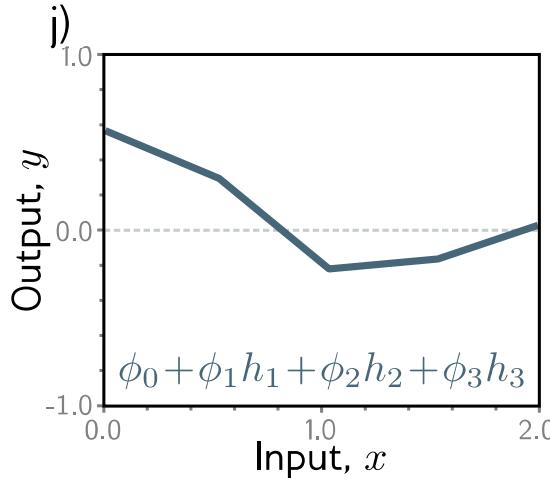
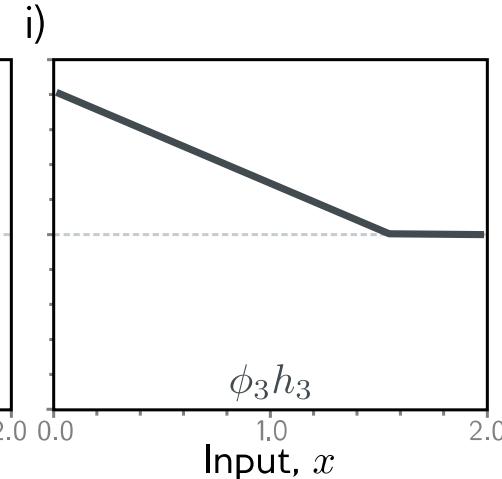
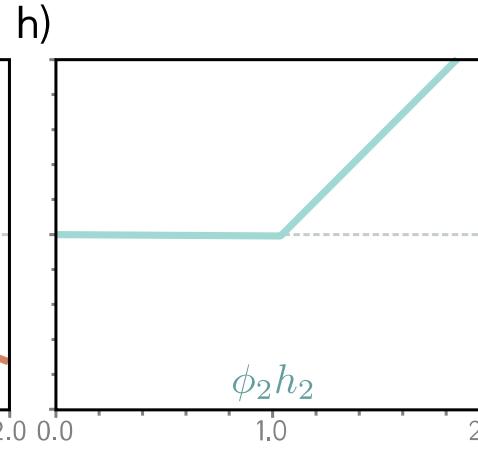
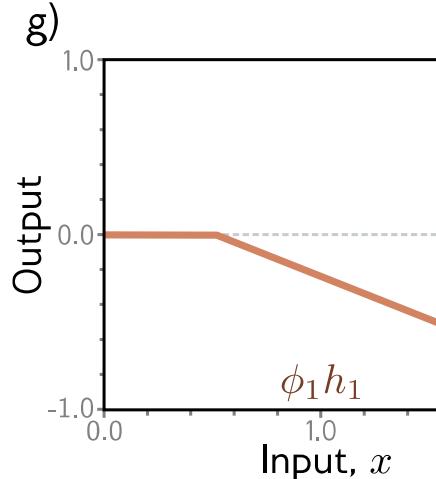
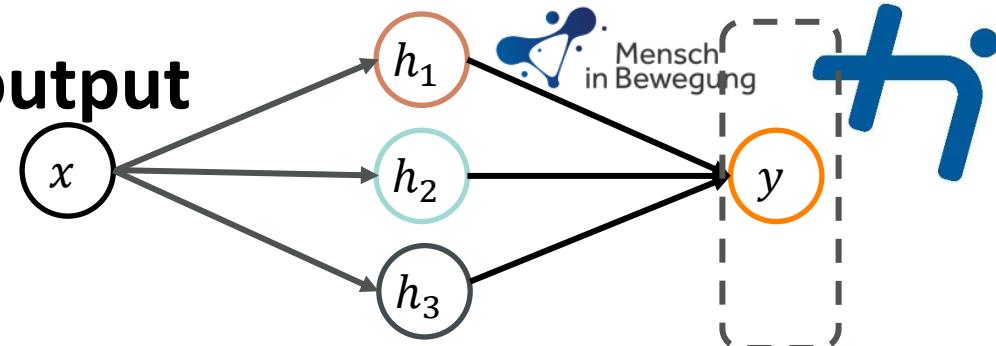
3. Weight the hidden units



4. Sum the weighted hidden units to create output

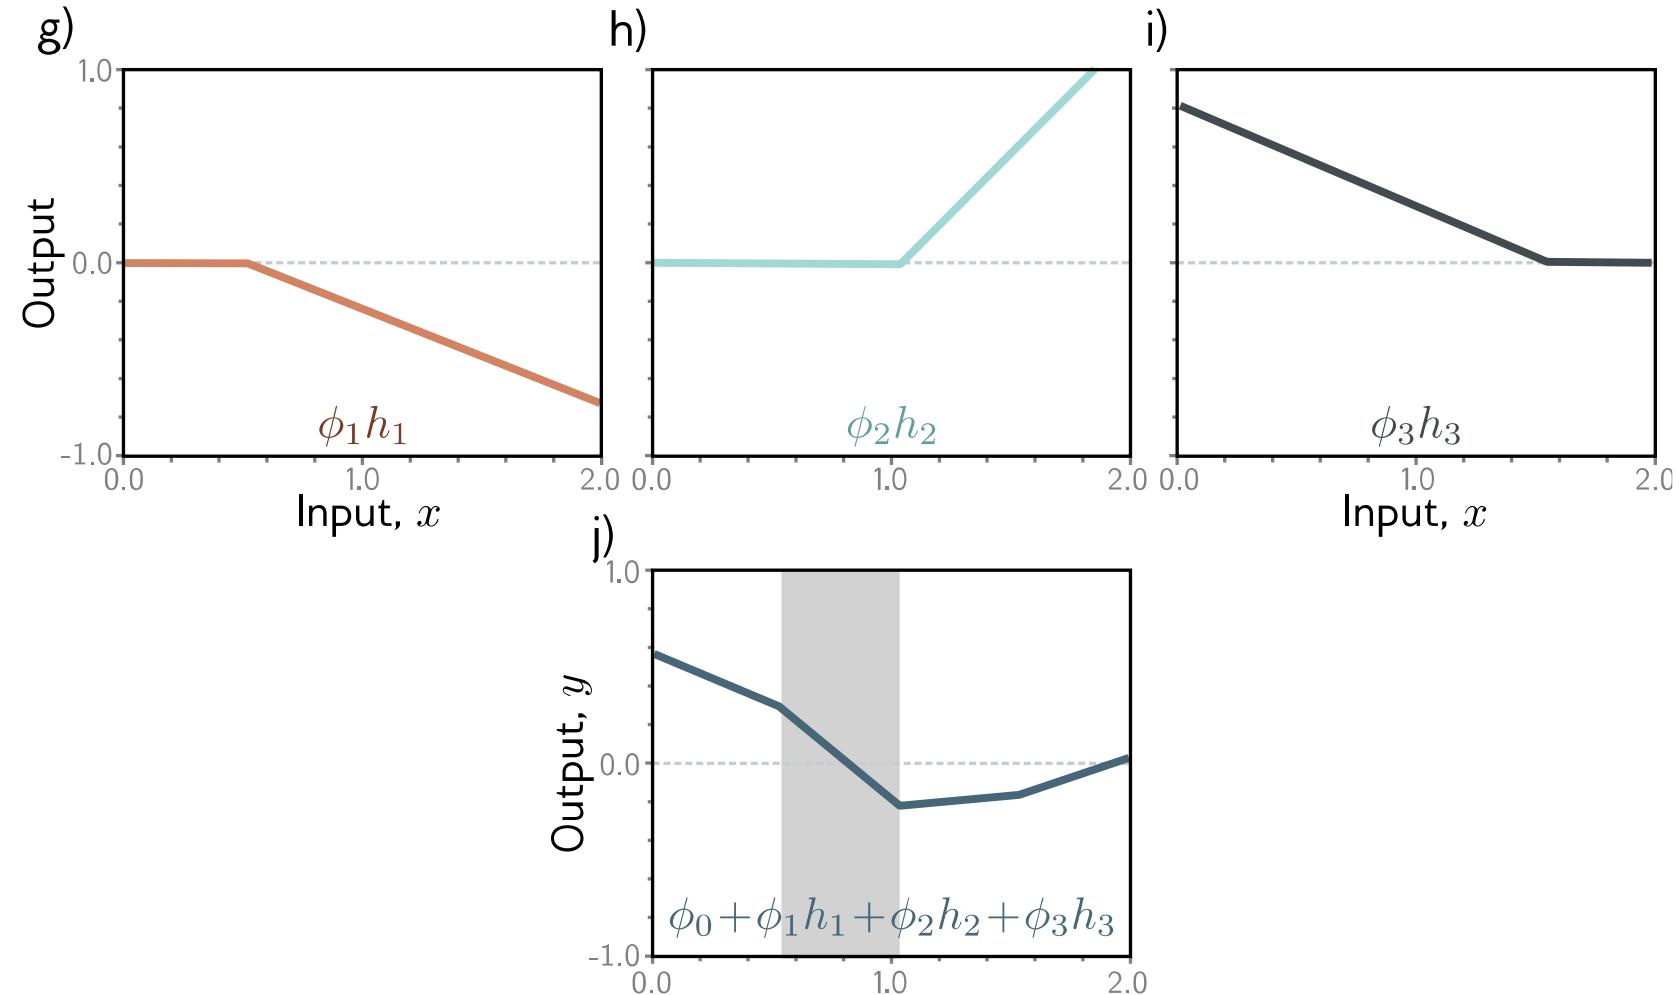


$$y = \phi_0 + \phi_1 h_1 + \phi_2 h_2 + \phi_3 h_3$$



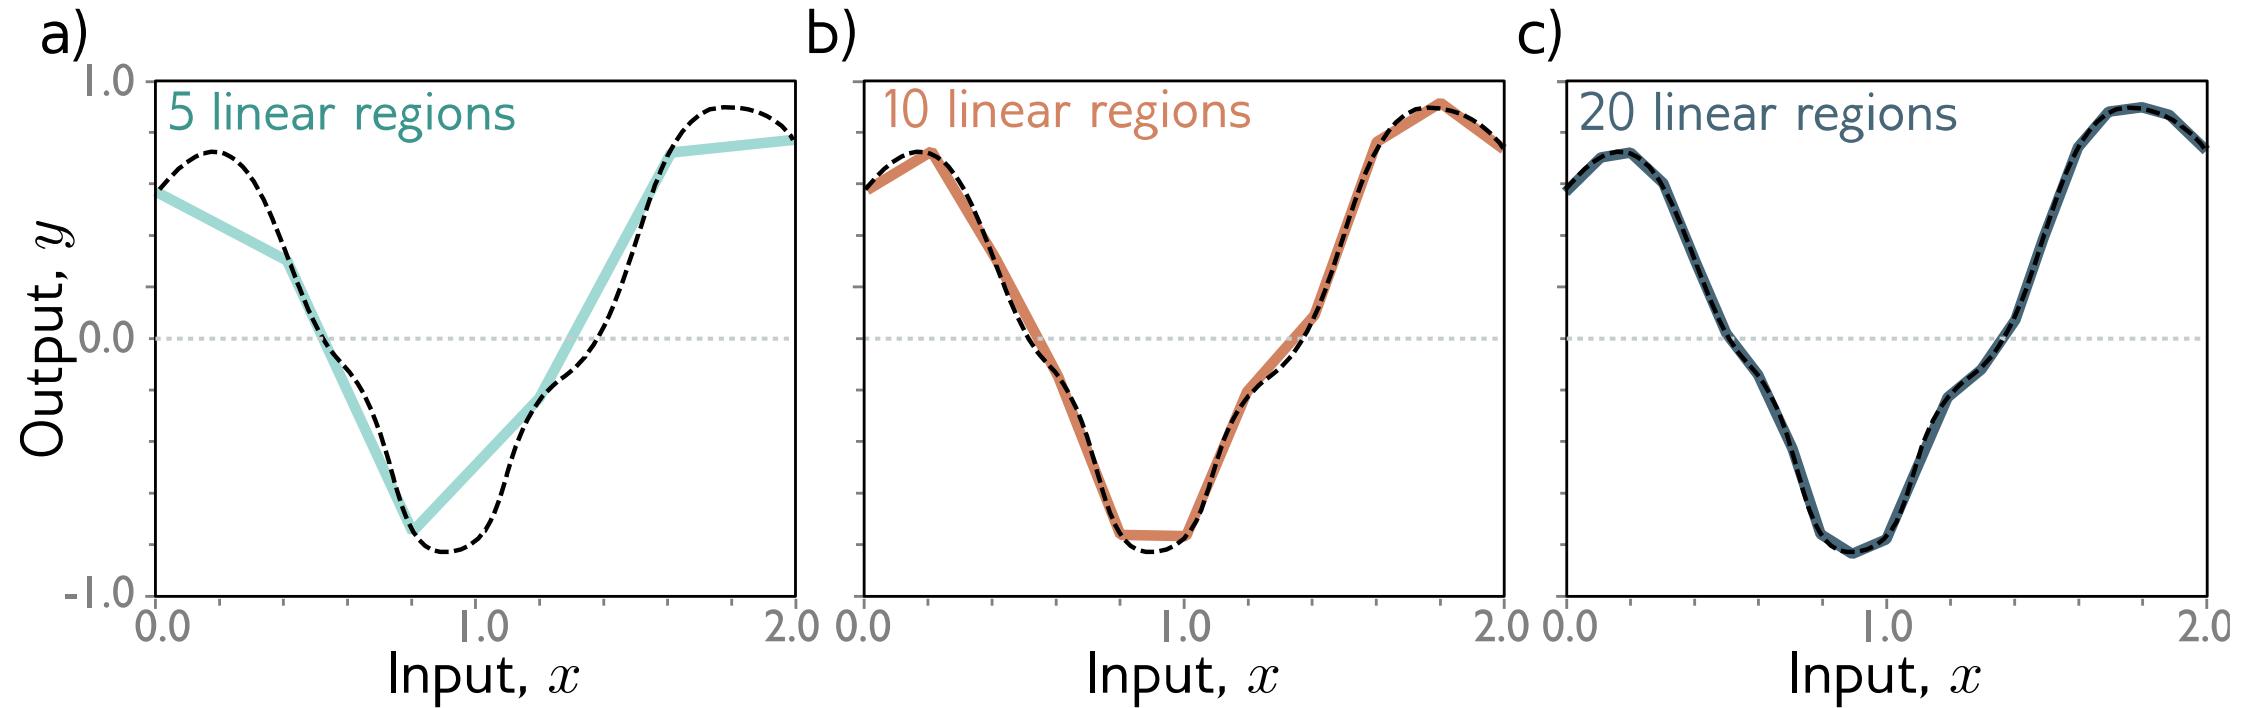
Activation pattern = which hidden units are activated

- Shaded region:
- Unit 1 **active**
- Unit 2 **inactive**
- Unit 3 **active**



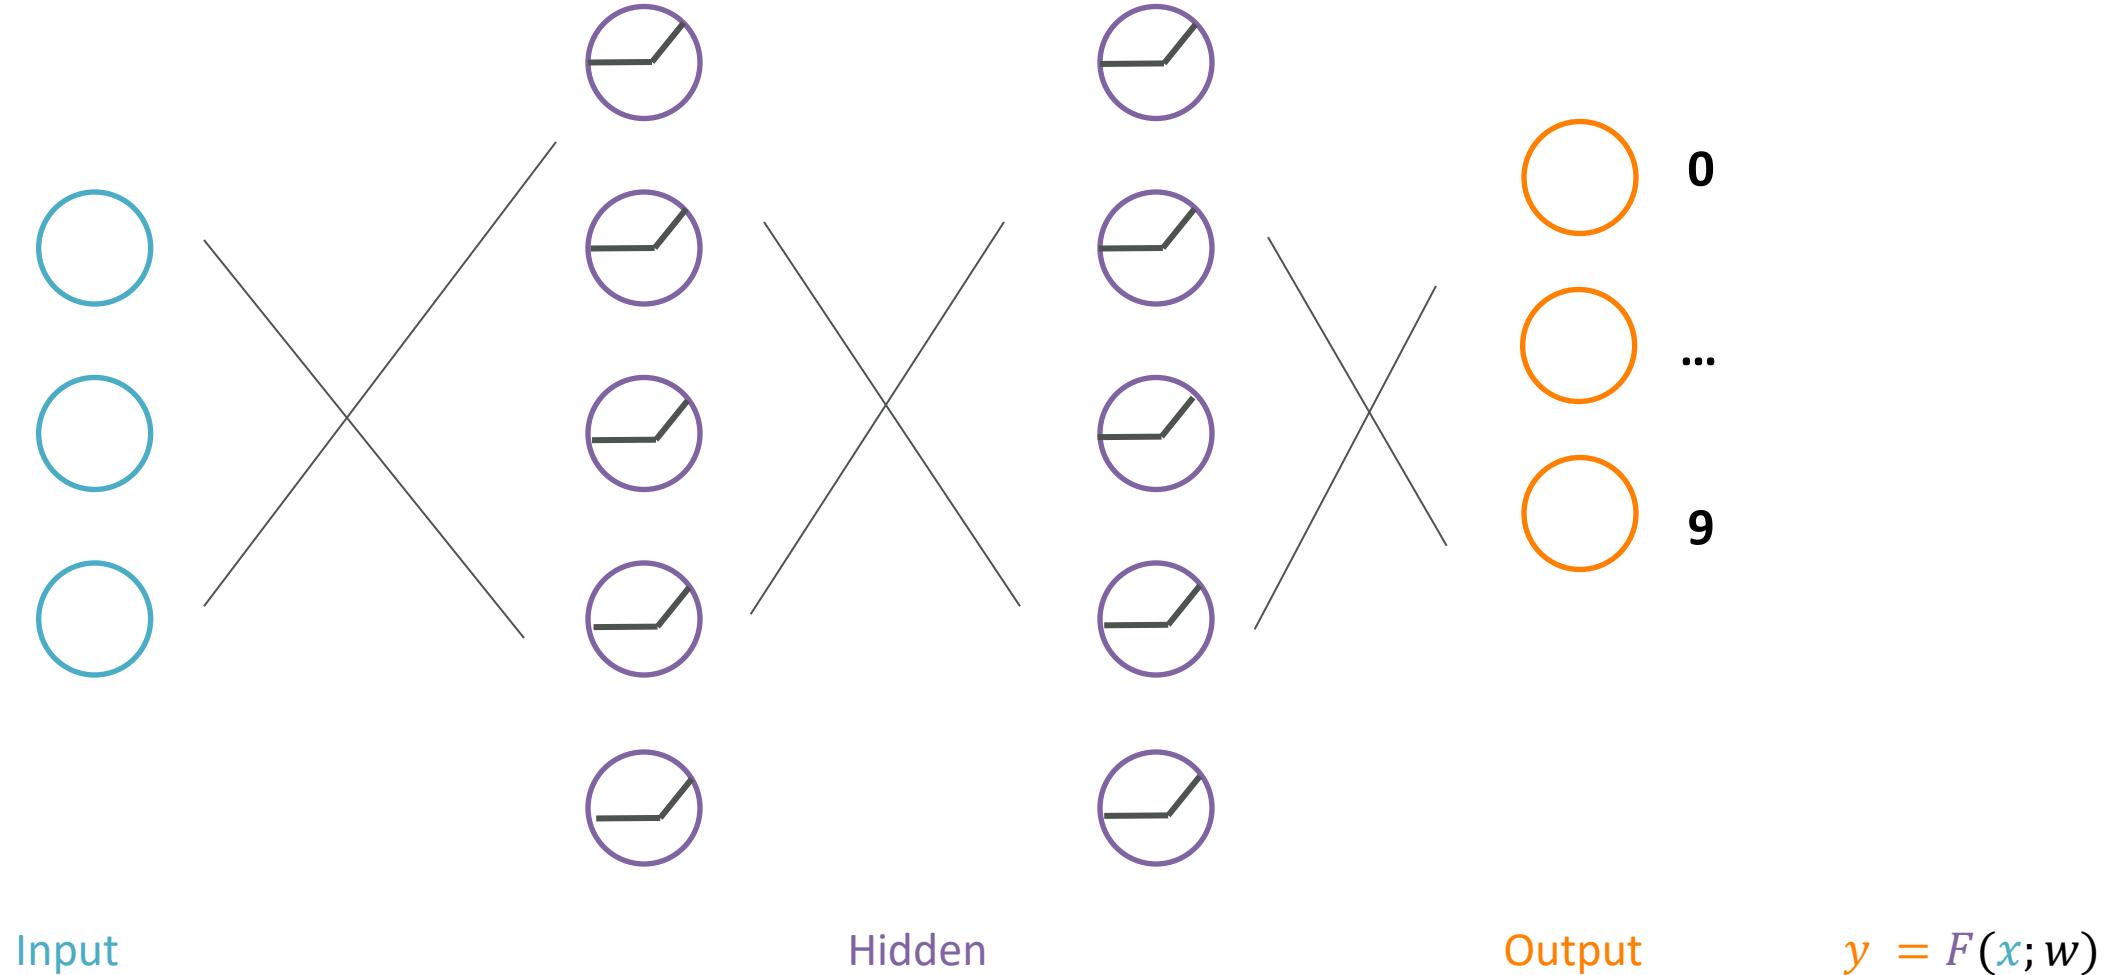
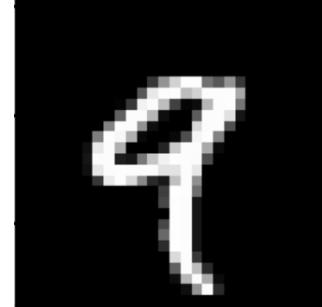
With enough hidden units...

... we can describe any 1D function to arbitrary accuracy

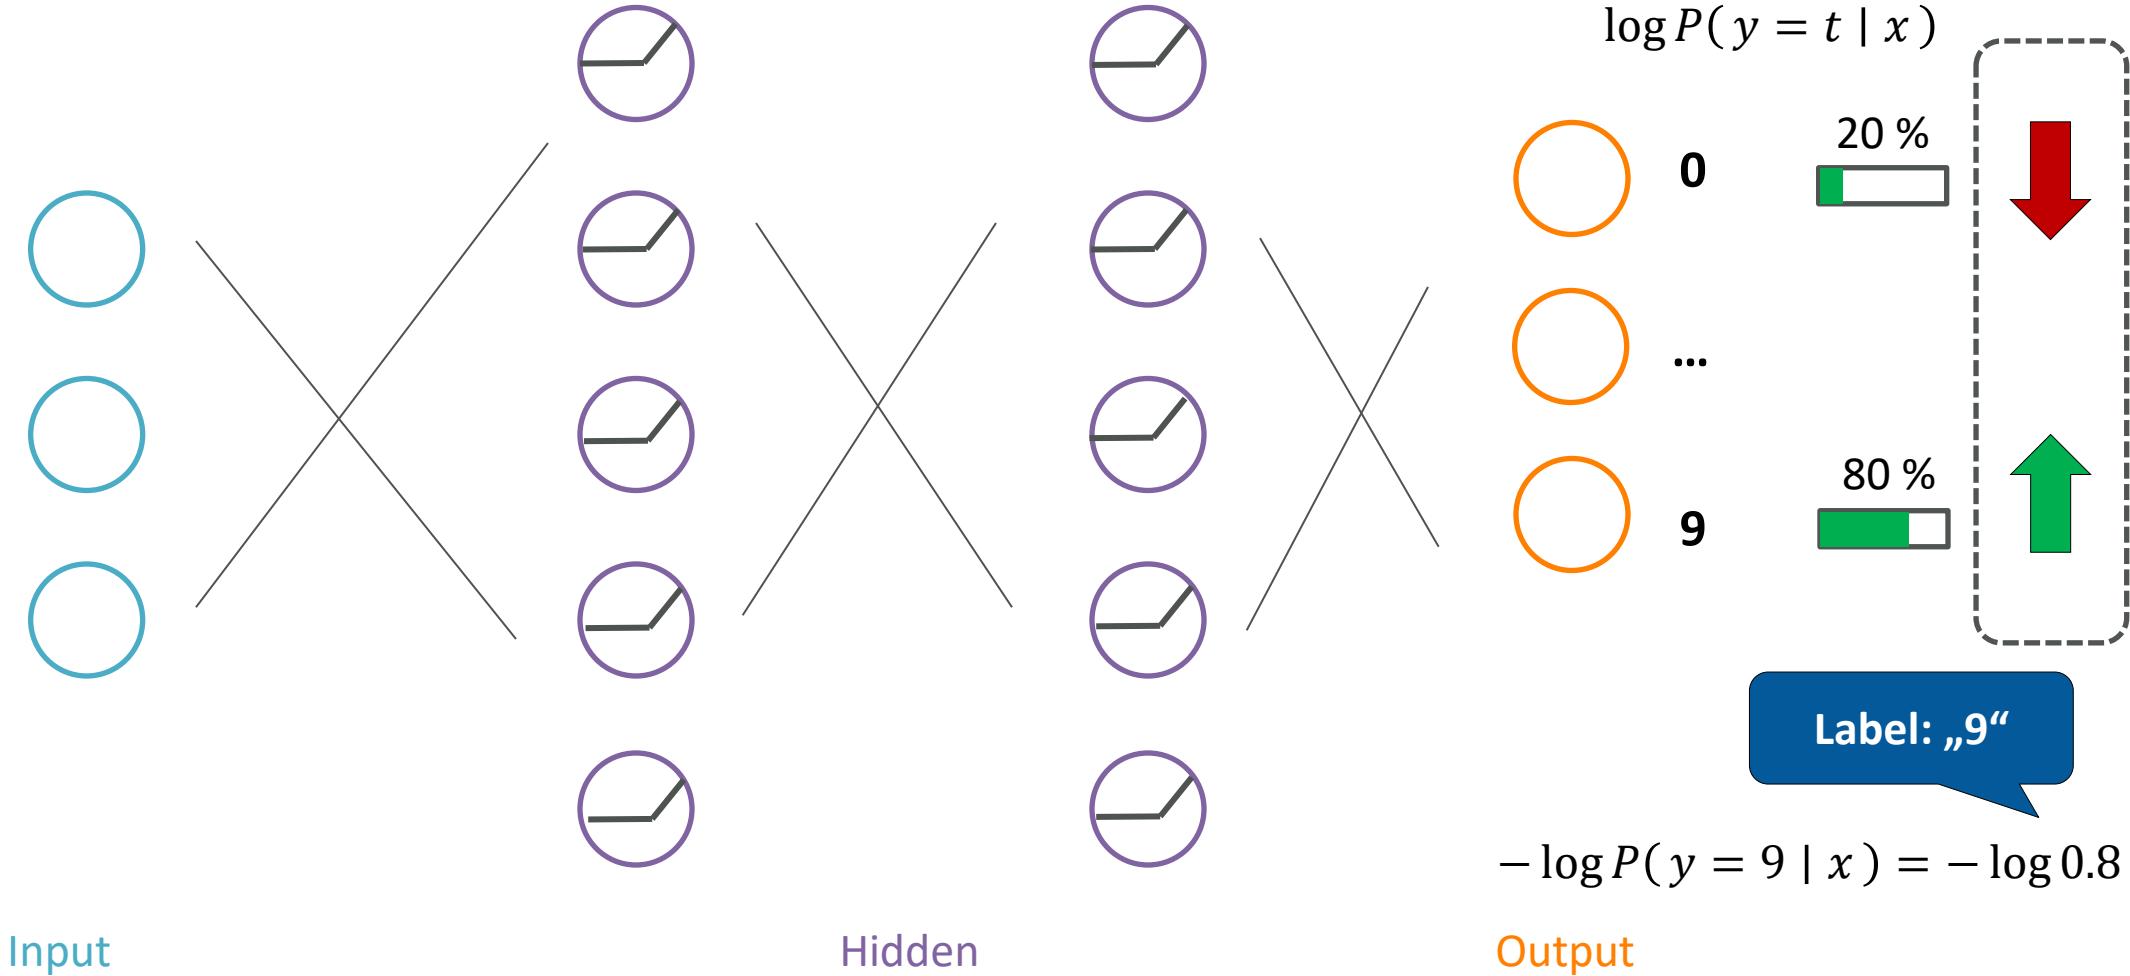
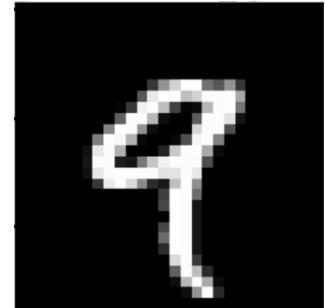


Automatically Training a Neural Network

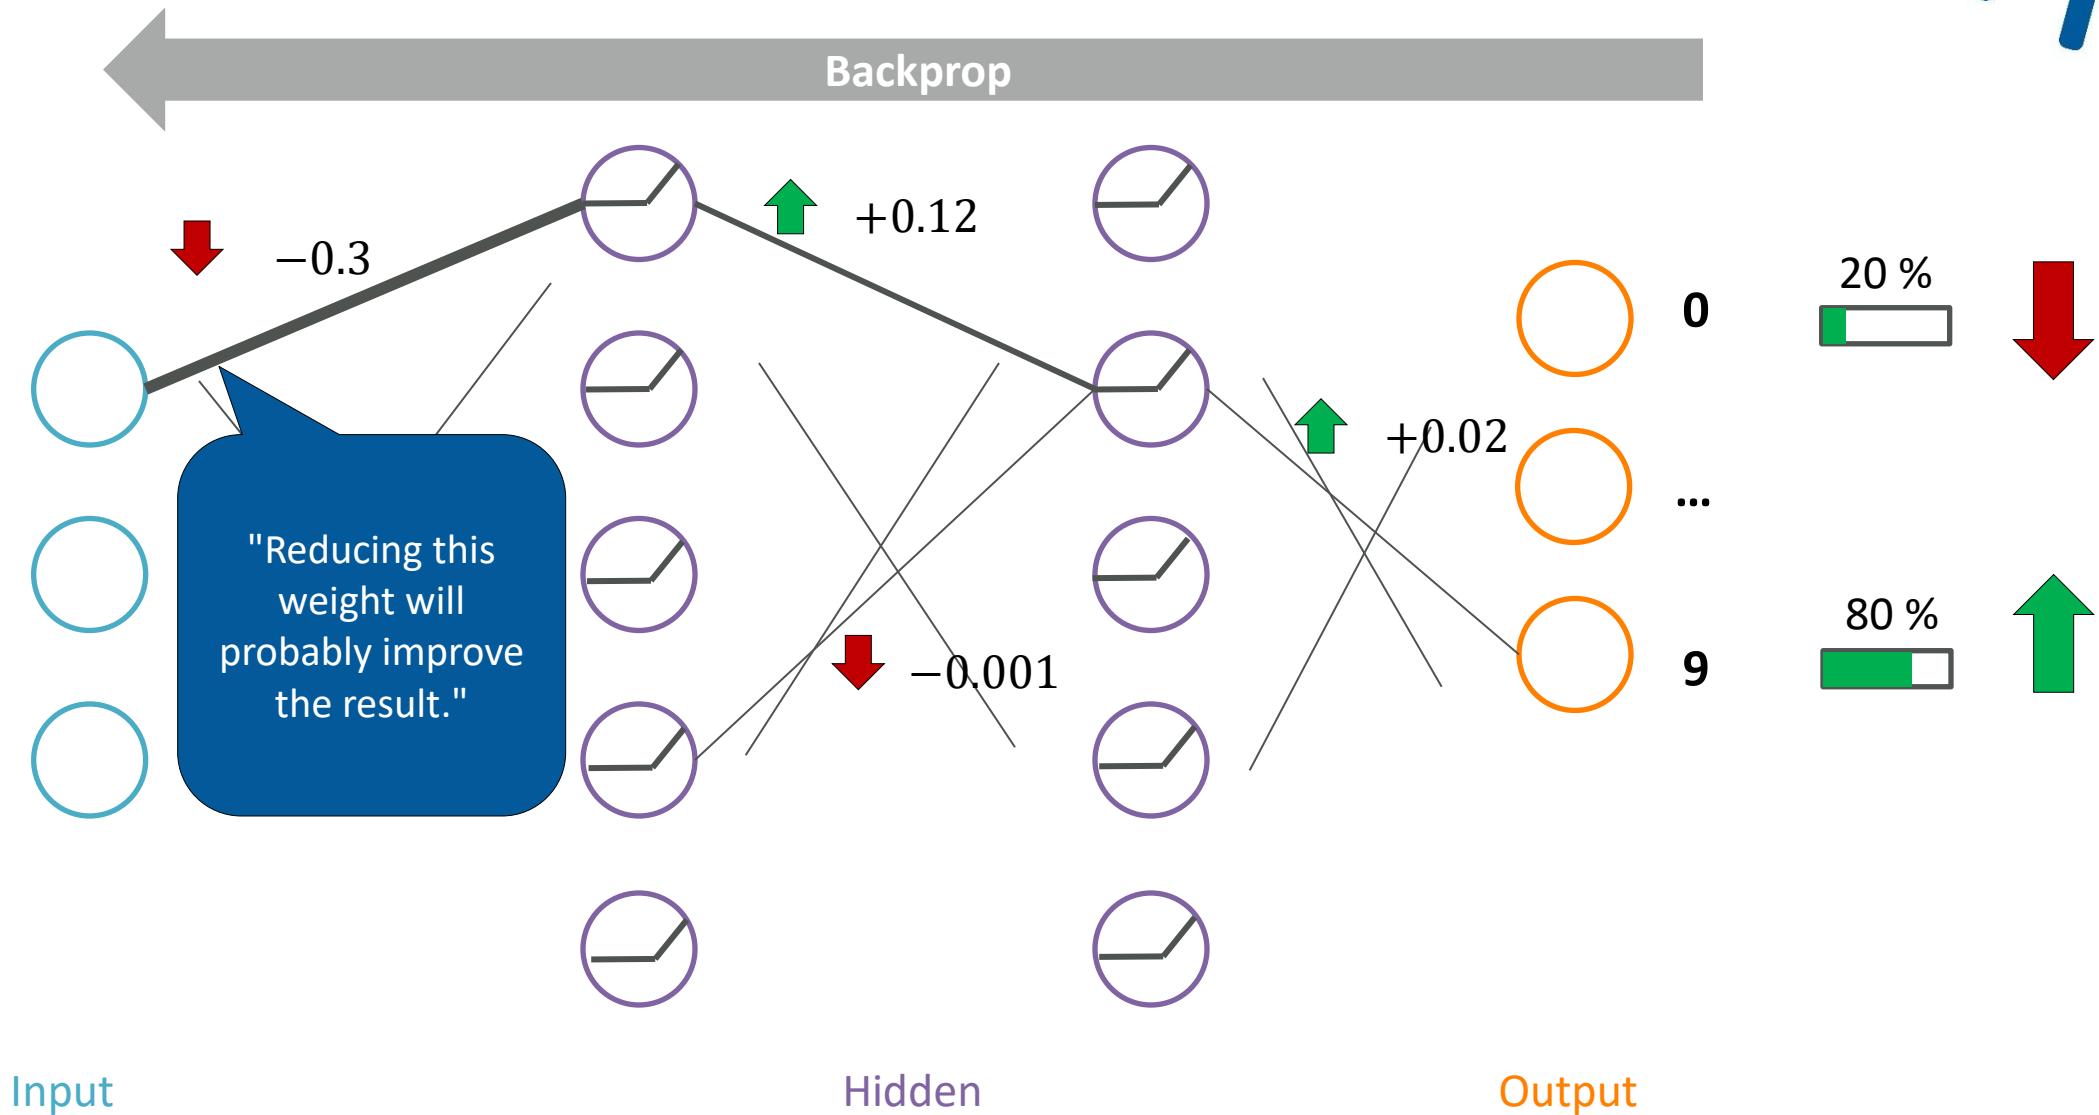
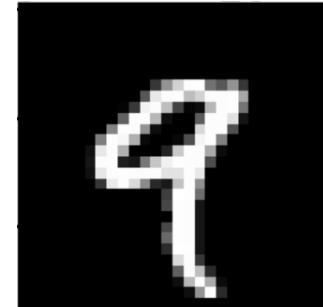
Training neural nets with backpropagation



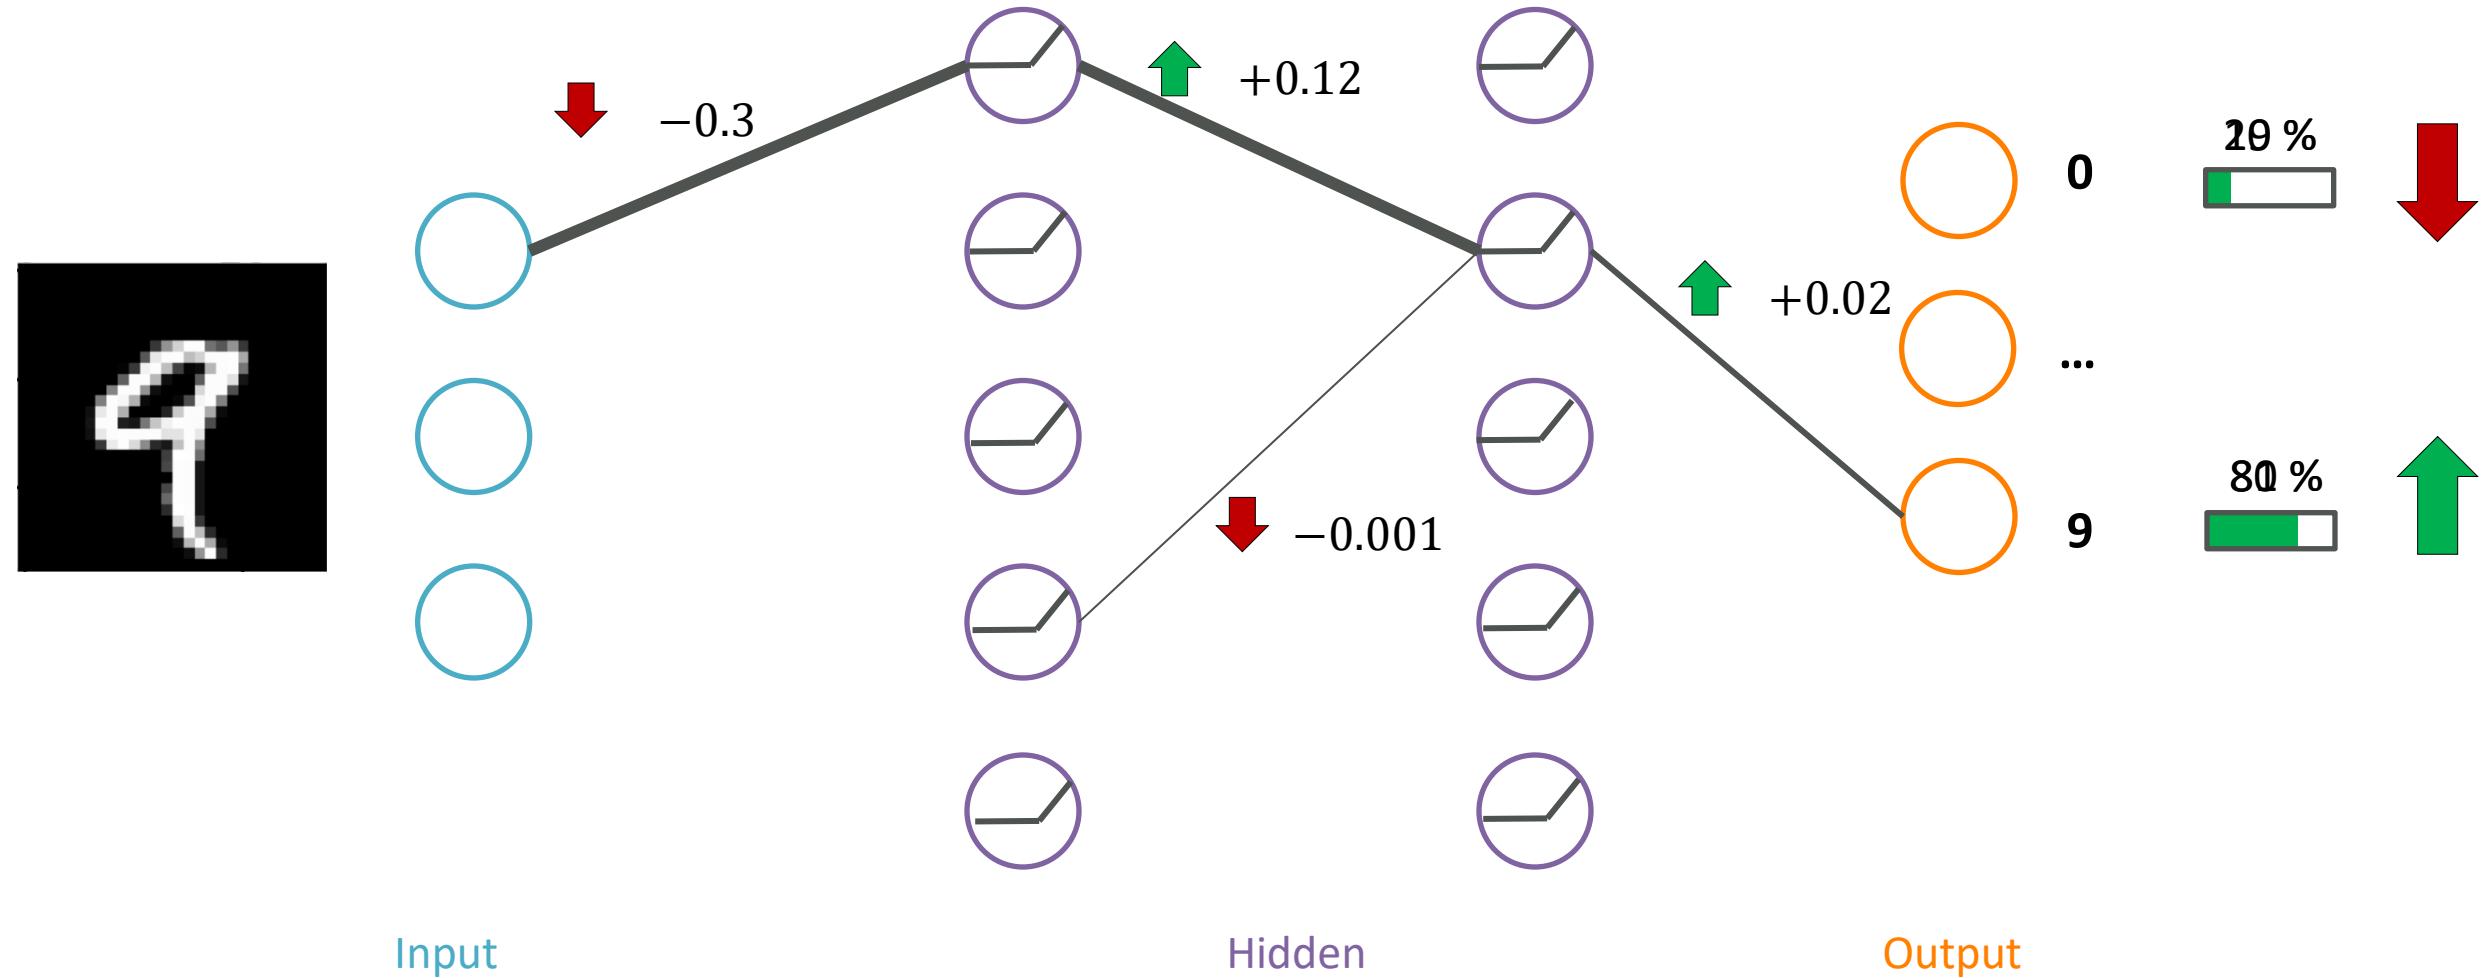
Forward Pass



Backpropagation – Gradient $\nabla_w L$



Backpropagation – Gradient update



Convolutional Neural Networks

Guiding Neuroscientific Principles

Hubel & Wiesel

Nobel Prize in Physiology or Medicine in 1981

1959

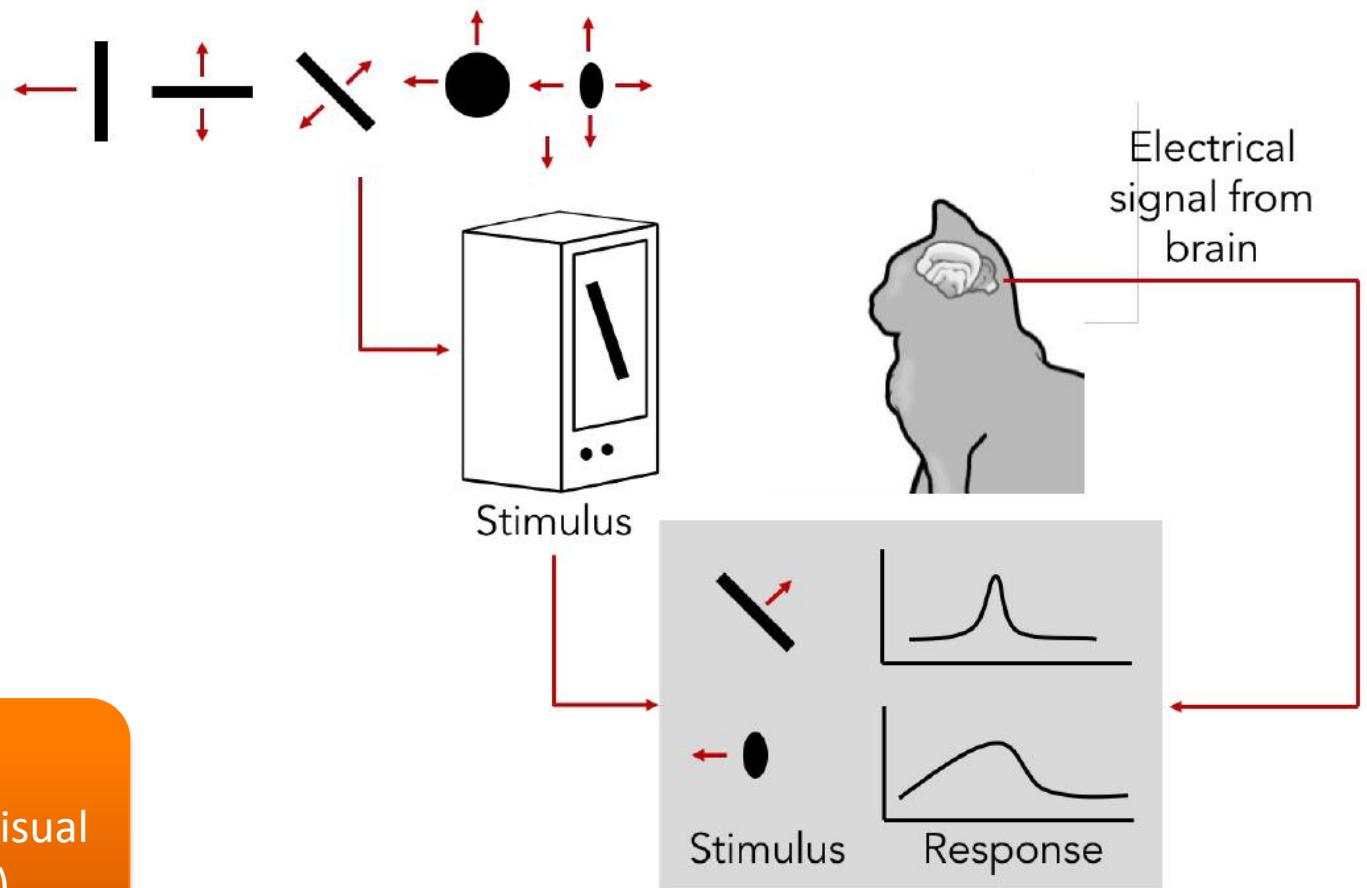
RECEPTIVE FIELDS OF SINGLE NEURONES IN
THE CAT'S STRIATE CORTEX

1962

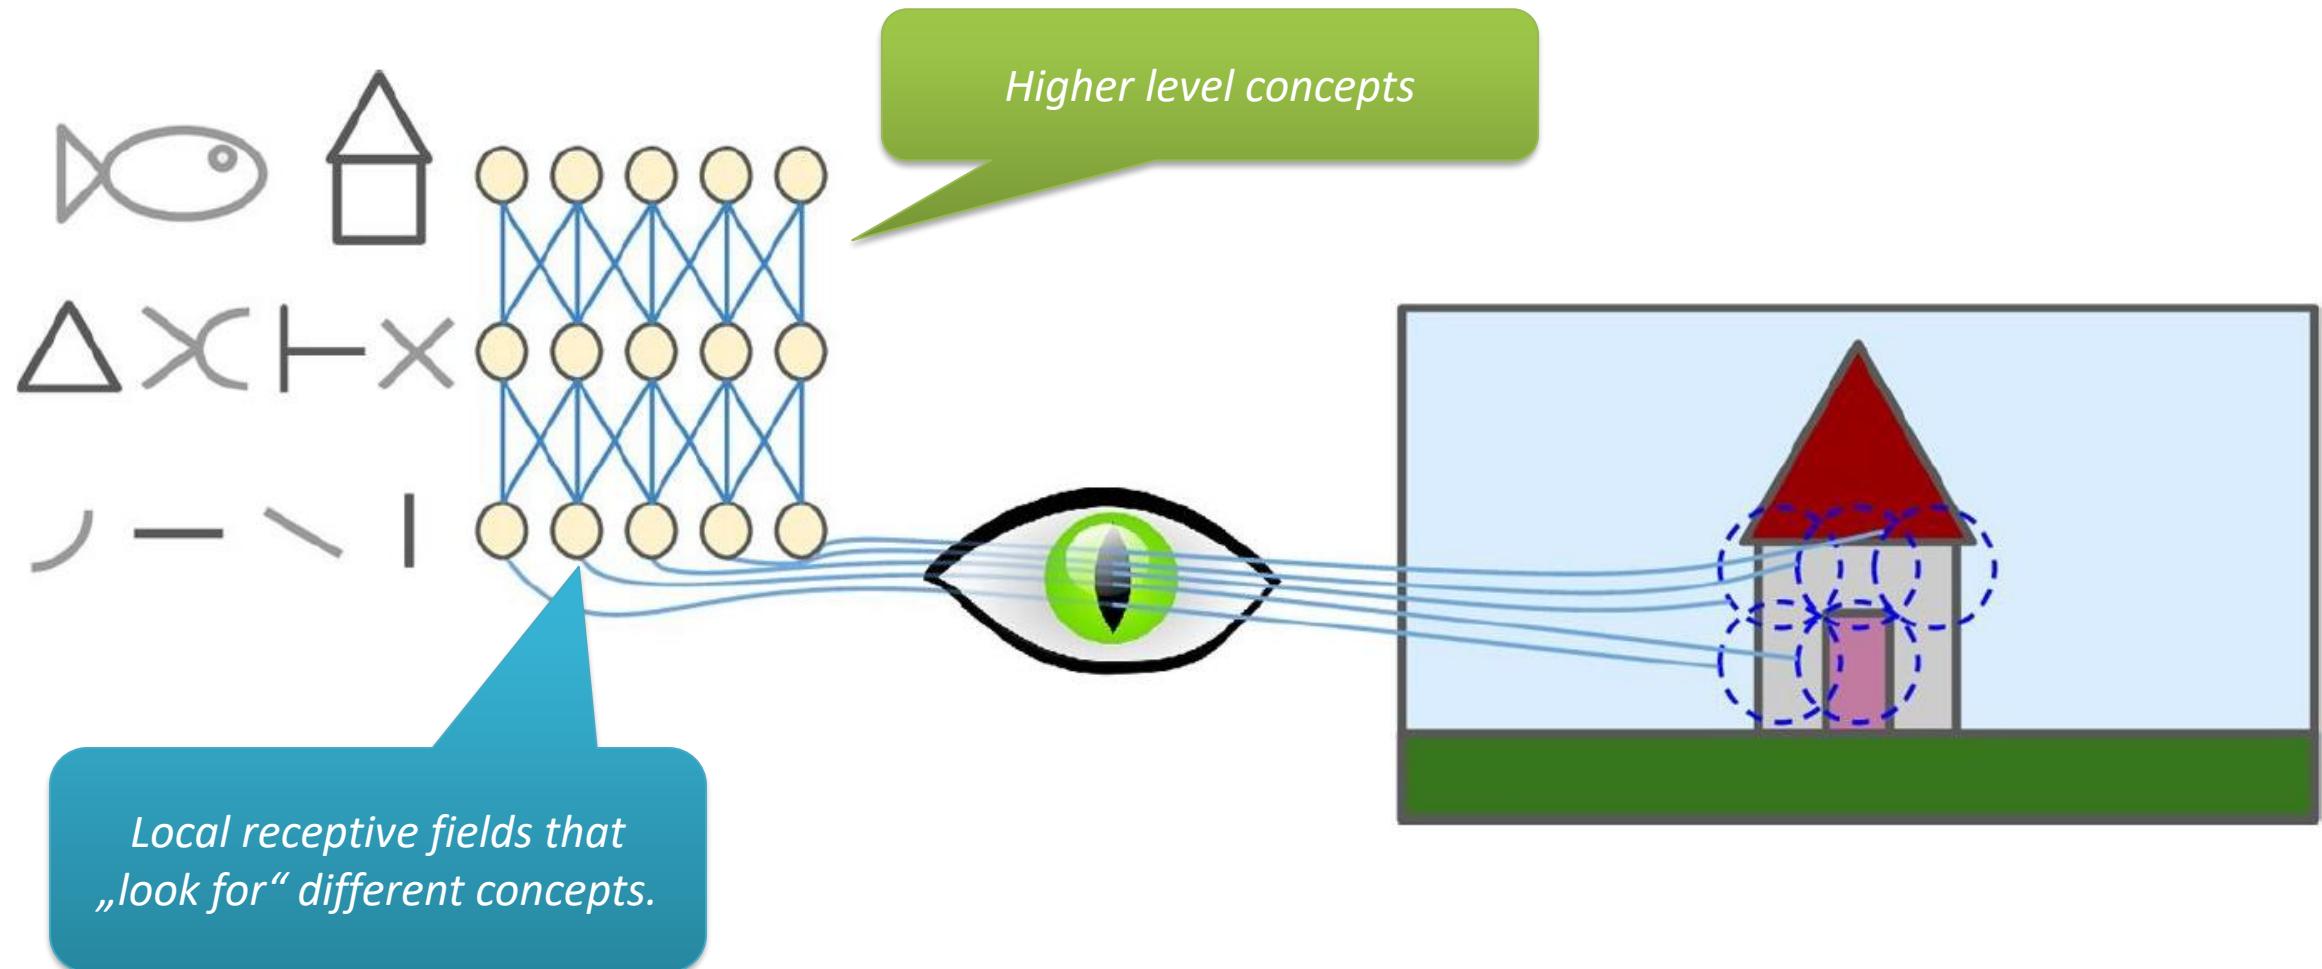
RECEPTIVE FIELDS, BINOCULAR INTERACTION
AND FUNCTIONAL ARCHITECTURE IN
THE CAT'S VISUAL CORTEX

Concept neurons:

Some neurons only react to certain patterns in the visual field (e.g., „horizontal lines“, „vertical lines“, etc.)



A Cartoon Impression Of Our Visual Cortex



[Hands-On Machine Learning with Scikit-Learn and Tensorflow, Géron, 2017]

Convolutions

Input Image In Channels



Original image (RGB)



R channel



G channel



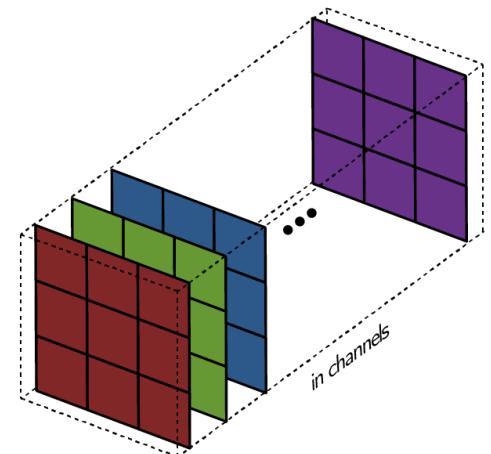
B channel

An image is a tensor with dimensions: $\mathbb{R}^{w \times h \times d}$

w(idth)

h(eight)

d(epth)

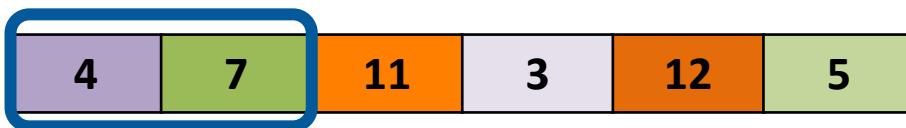


Discrete Convolutions as feature extractors

Filter kernel

-1	1
----	---

-1 1



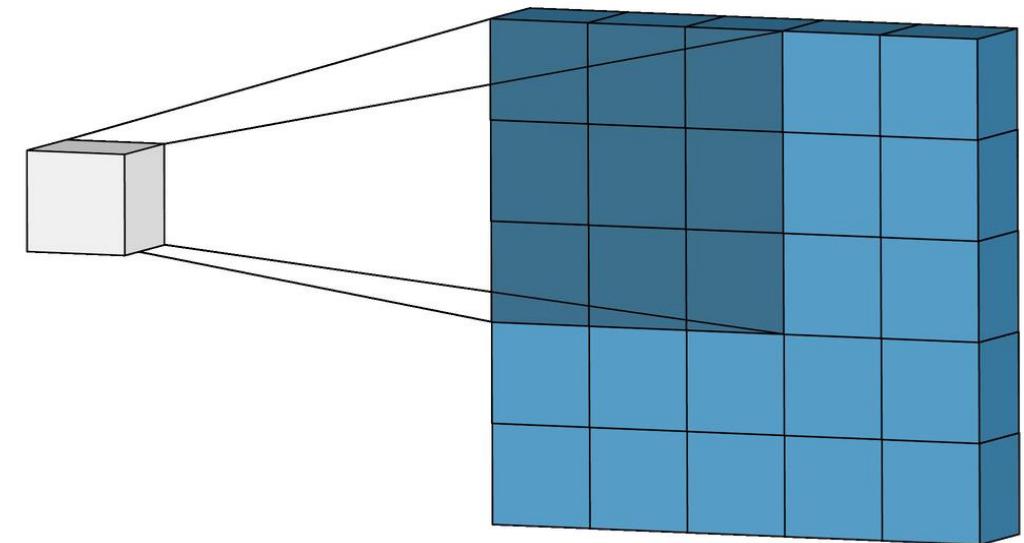
3	4	-8	9	-7
---	---	----	---	----

Convolutions: Core Idea

- Exploit **grid-like spatial structure** (1D time series, 2D images, 3D videos, etc.)
- Restrict input locations to neighborhoods (receptive fields)

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0



[<https://towardsdatascience.com/intuitively-understanding-convolutions-for-deep-learning-1f6f42faee1>]

Convolution

0.4	0.4	0.6	0	0	0	0	0
0	0	0.5	0	0	0	0	0
0	0	0.6	0	0	0	0	0
0	0	0.6	0	0	0	0	0
0	0	0.4	0	0	0	0	0
0	0	0.2	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

Convolve with

1	0	-1
1	0	-1
1	0	-1

-1.3	0.4	1.7	0	0	0
-1.7	0	1.7	0	0	0
-1.6	0	1.6	0	0	0
-1.2	0	1.2	0	0	0
-0.6	0	0.6	0	0	0
-0.2	0	0.2	0	0	0

$$0.4 - 0.6 - 0.5 - 0.6 = -1.3$$

$$0.6 + 0.5 + 0.6 = 1.7$$

Convolution „preserves“ shift operations

0.4	0.4	0.6	0	0	0	0	0
0	0	0.5	0	0	0	0	0
0	0	0.6	0	0	0	0	0
0	0	0.6	0	0	0	0	0
0	0	0.4	0	0	0	0	0
0	0	0.2	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$$\begin{matrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{matrix}$$

-1.3	0.4	1.7	0	0	0
-1.7	0	1.7	0	0	0
-1.6	0	1.6	0	0	0
-1.2	0	1.2	0	0	0
-0.6	0	0.6	0	0	0
-0.2	0	0.2	0	0	0

0	0	0.4	0.4	0.6	0	0	0
0	0	0	0	0.5	0	0	0
0	0	0	0	0.6	0	0	0
0	0	0	0	0.6	0	0	0
0	0	0	0	0.4	0	0	0
0	0	0	0	0.2	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$$\begin{matrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{matrix}$$

-0.4	-0.4	-1.3	0.4	1.7	0
0	0	-1.7	0	1.7	0
0	0	-1.6	0	1.6	0
0	0	-1.2	0	1.2	0
0	0	-0.6	0	0.6	0
0	0	-0.2	0	0.2	0

Convolution: A Sobel Filter For Edge Detection



I

$$I * G_x$$

$$* \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

↗

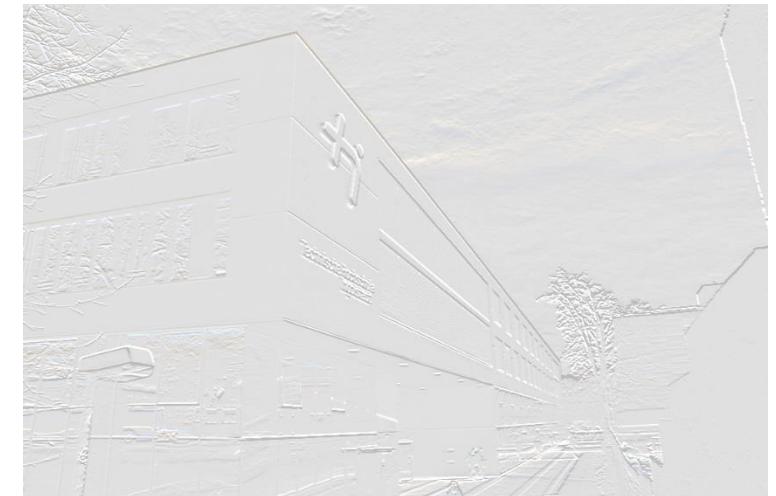
$$I * G_y$$

$$* \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

↘



$I * G_x$



$I * G_y$

Task



Play around with the excel sheet!

https://loma-aimotion.github.io/ml_visualizations/convolution_demo.xlsx

Convolution: A Sobel Filter For Edge Detection



I

$$I * G_x$$

$$* \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

↗

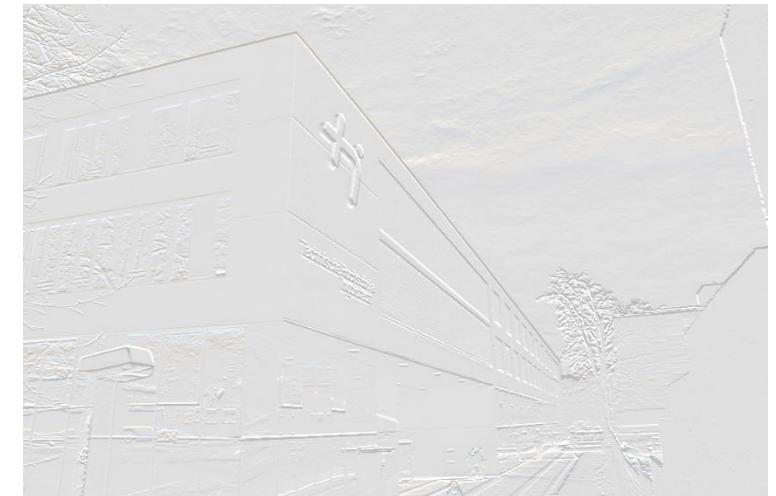
$$I * G_y$$

$$* \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

↘



$I * G_x$

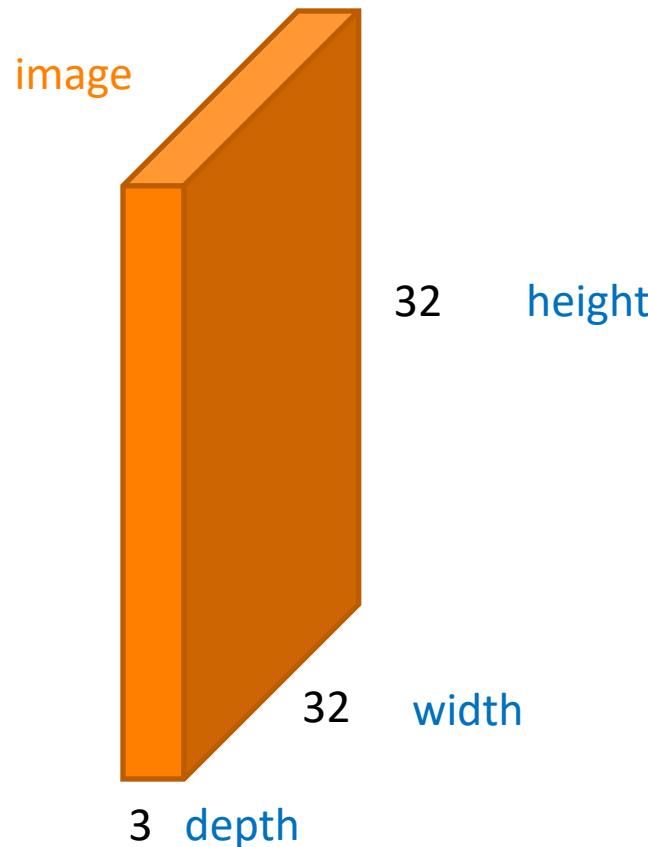


$I * G_y$

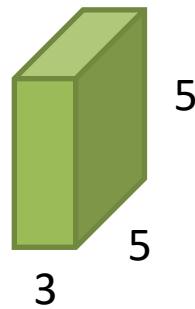
Convolutional Layers

Convolutional Layer

Preserve spatial structure of input instead of flattening



kernel: $5 \times 5 \times 3$

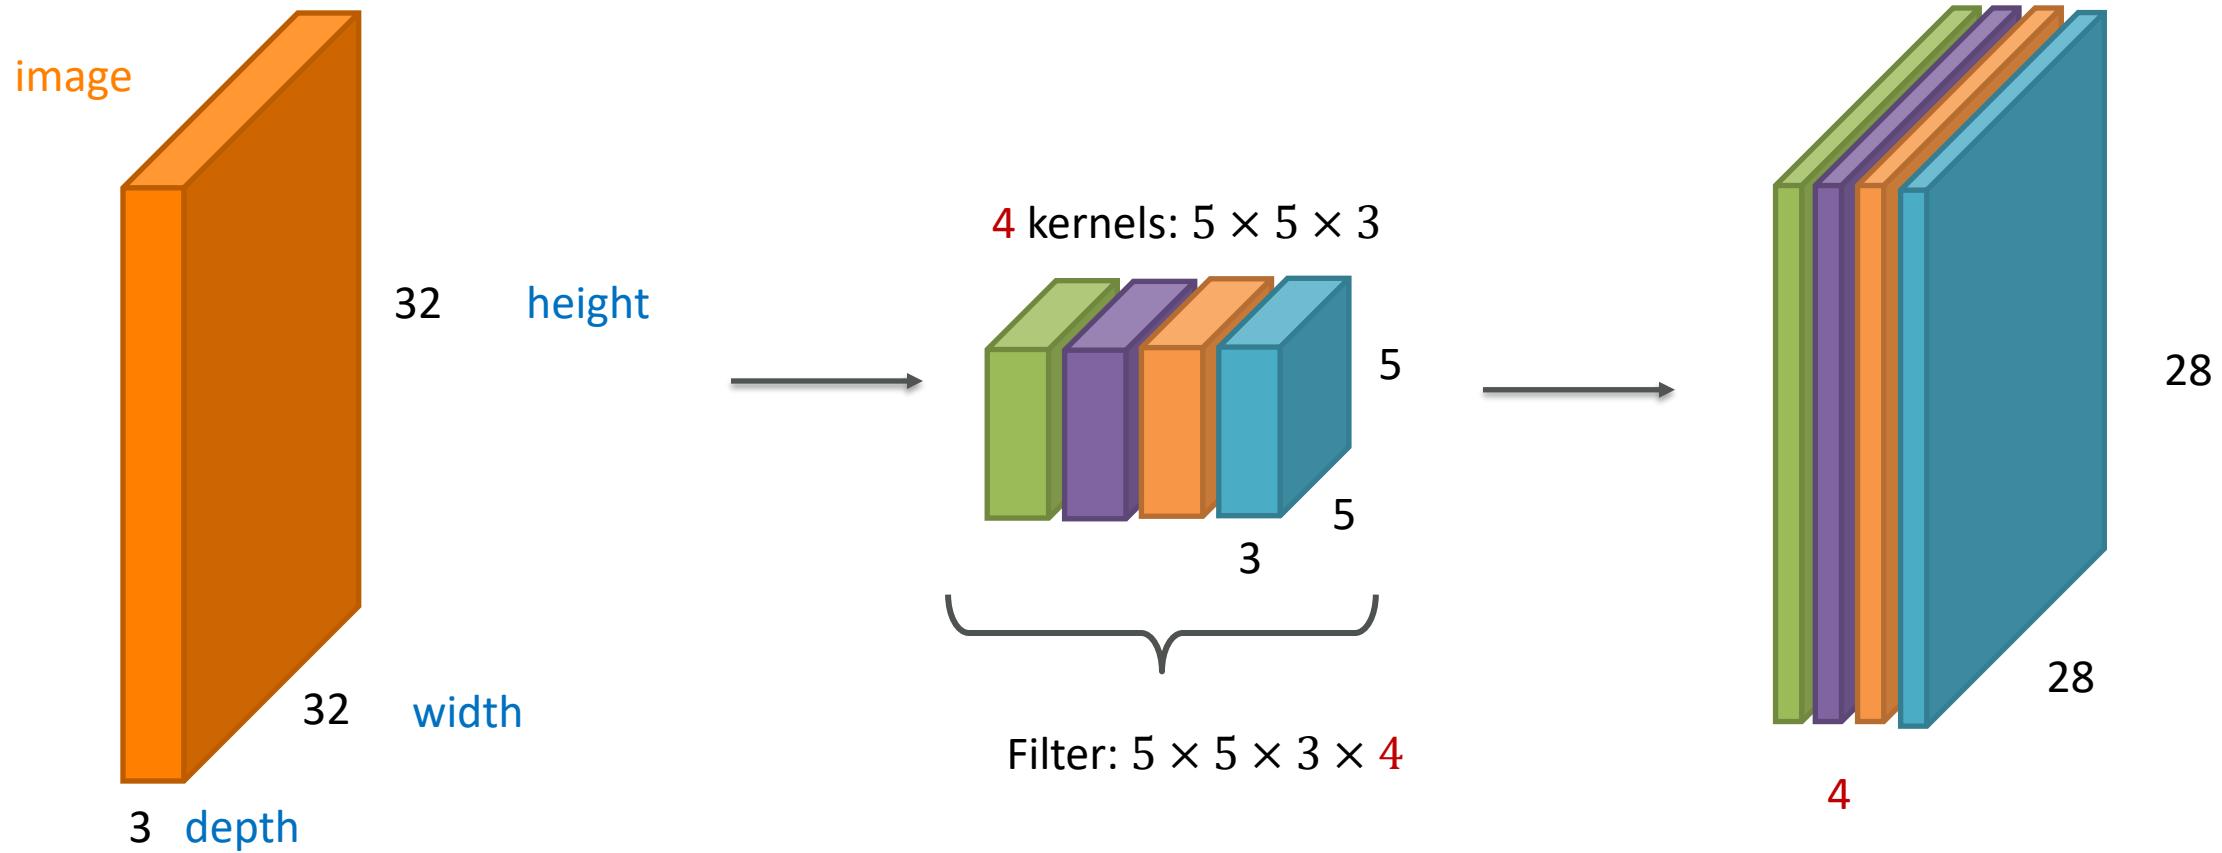


Convolution means „sliding“ the filter over the image and computing dot products.

Kernels typically extend through all channels ($d = 3$)!

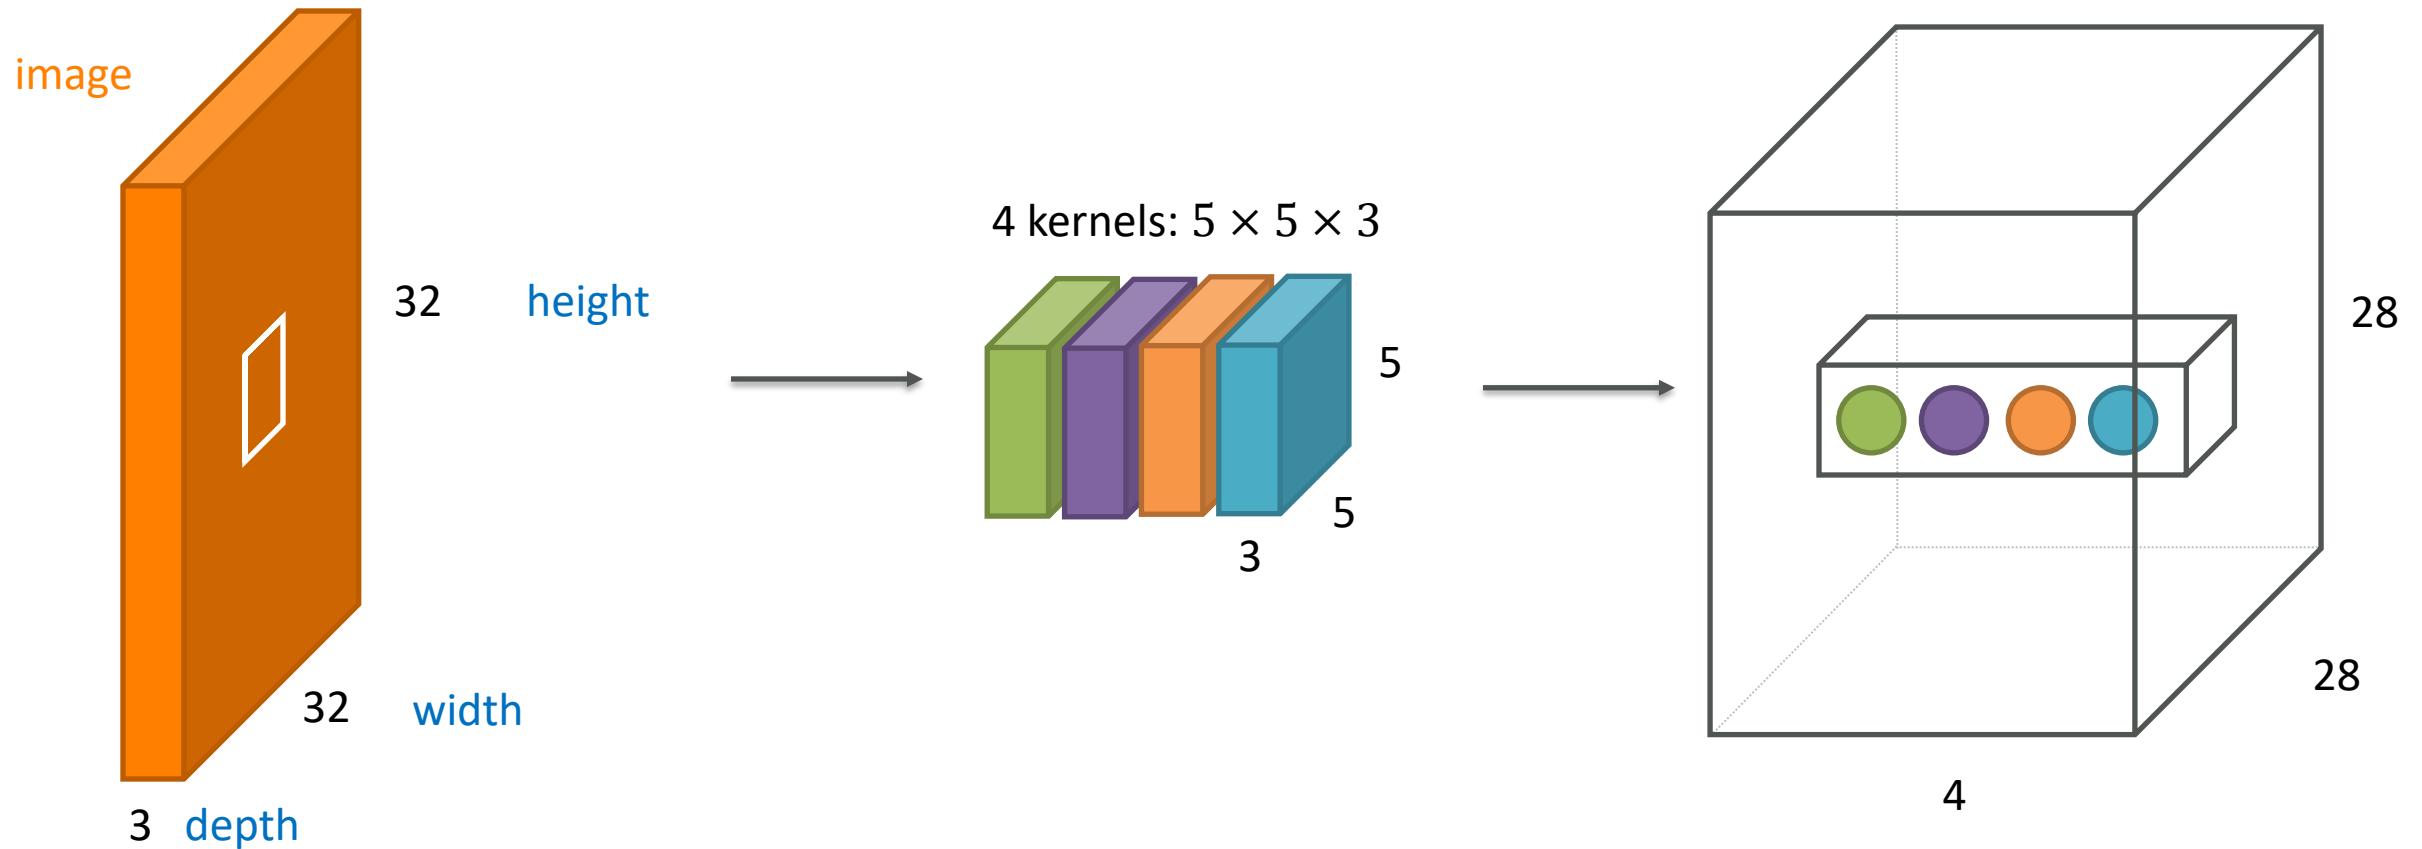
Convolutional Layer

Each kernel results in a different activation map

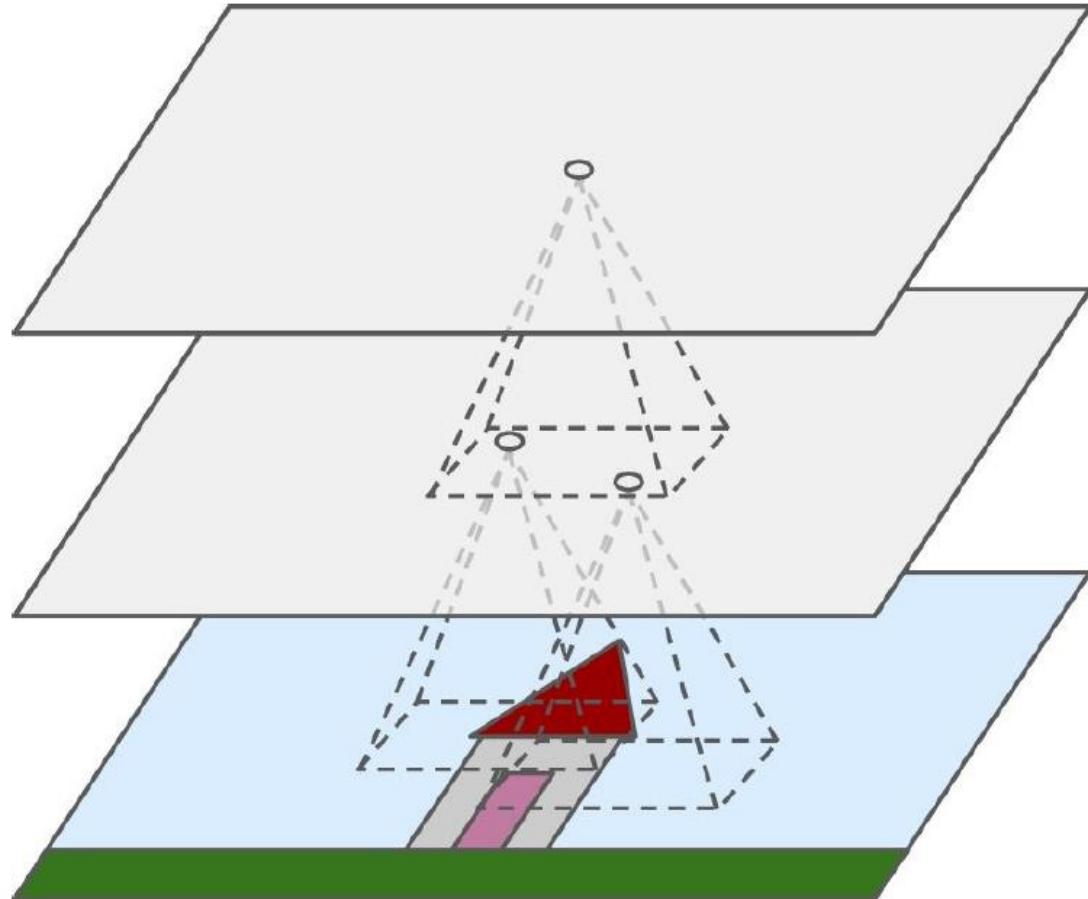


Convolutional Layer

Each kernel looks at the same region of the input, but for different things

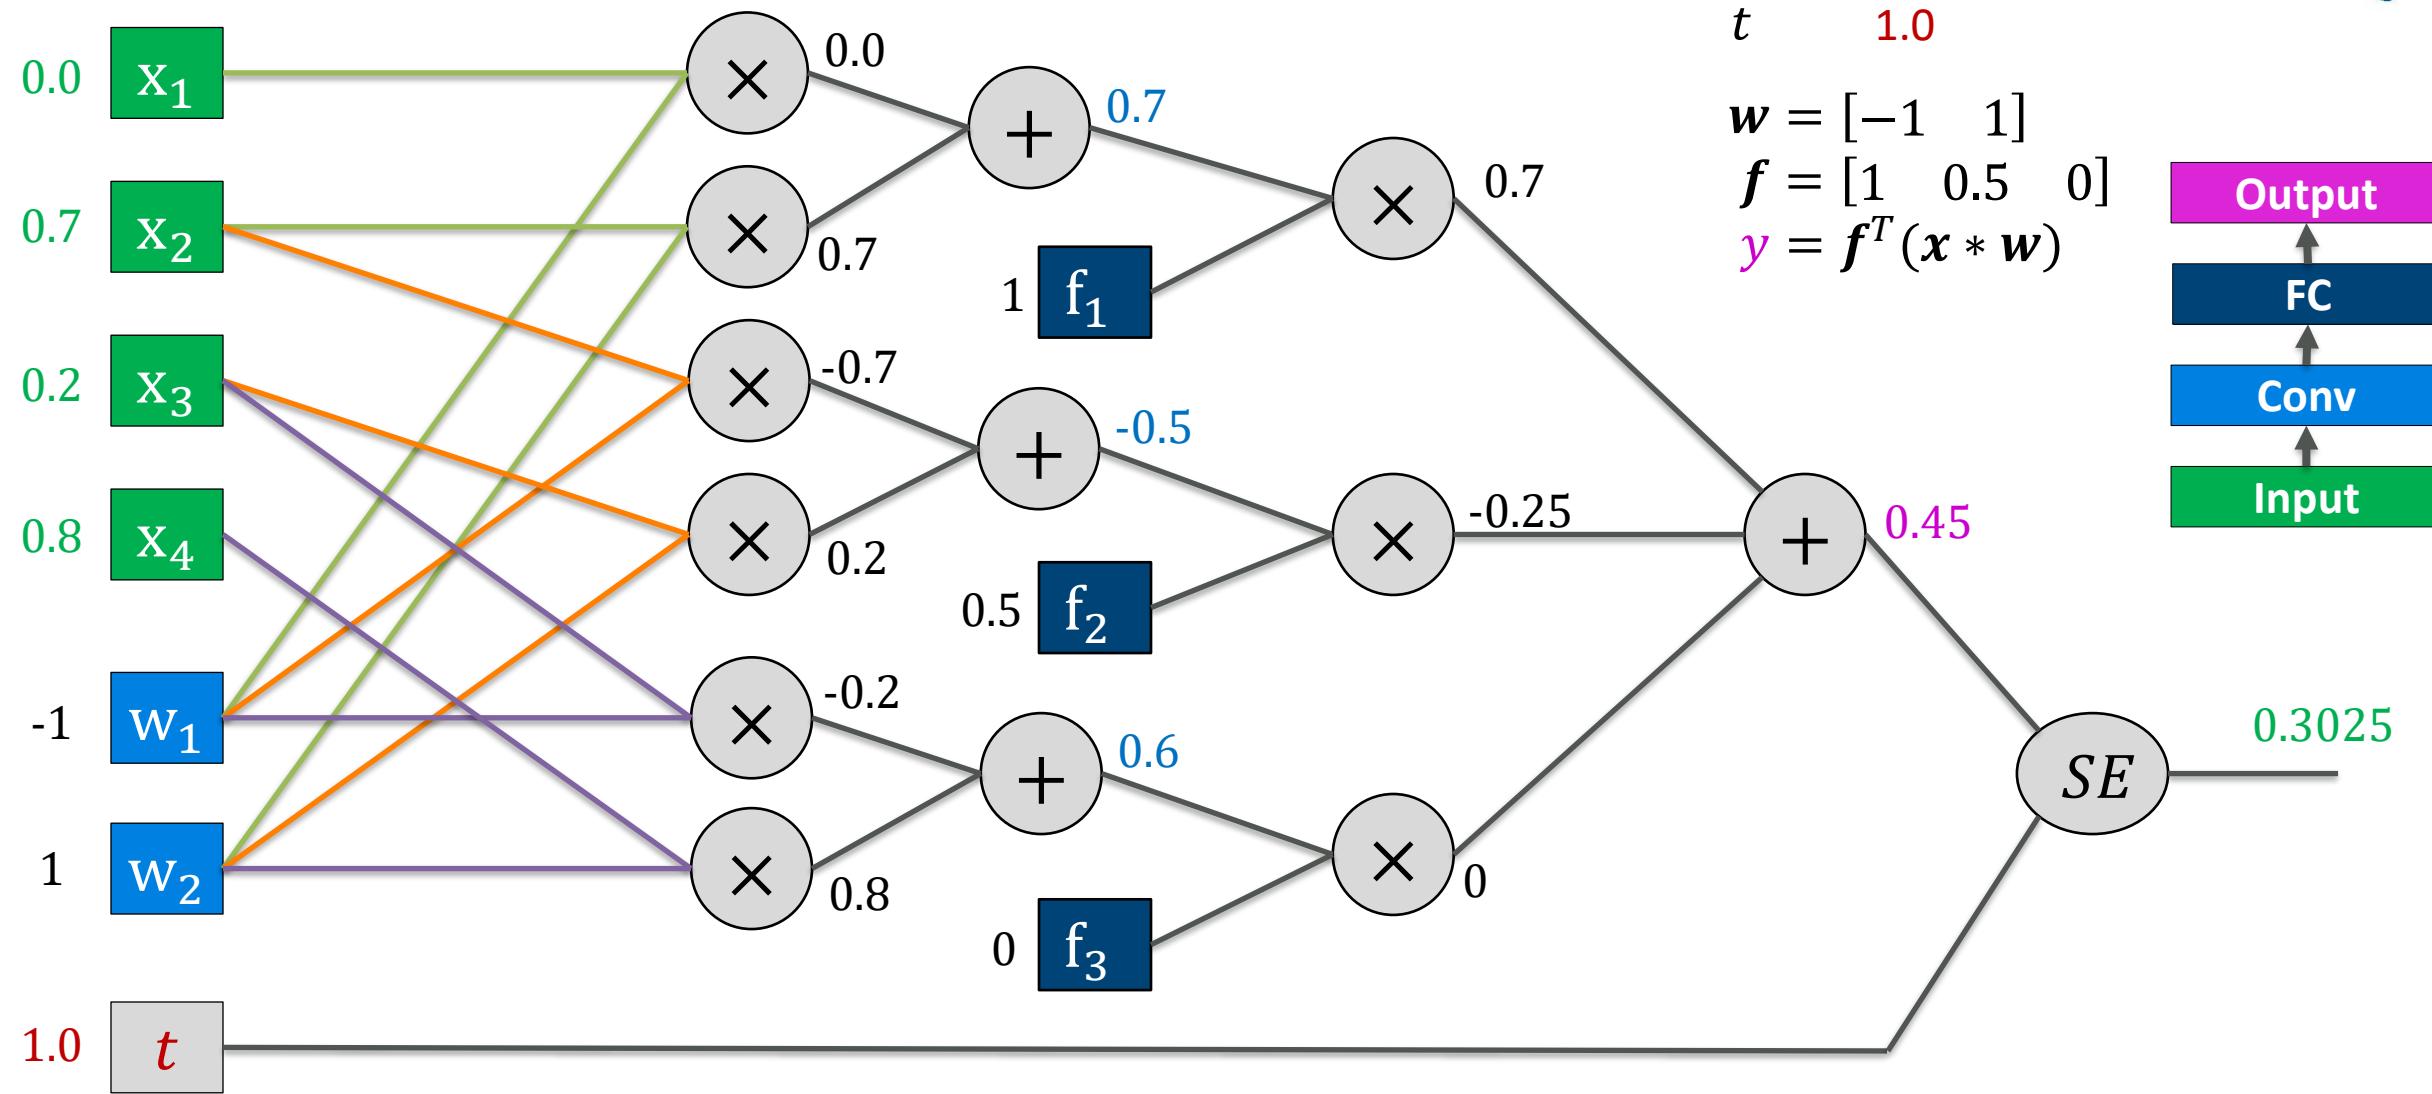


Interpretation of Convolutional Layers as Receptive Fields

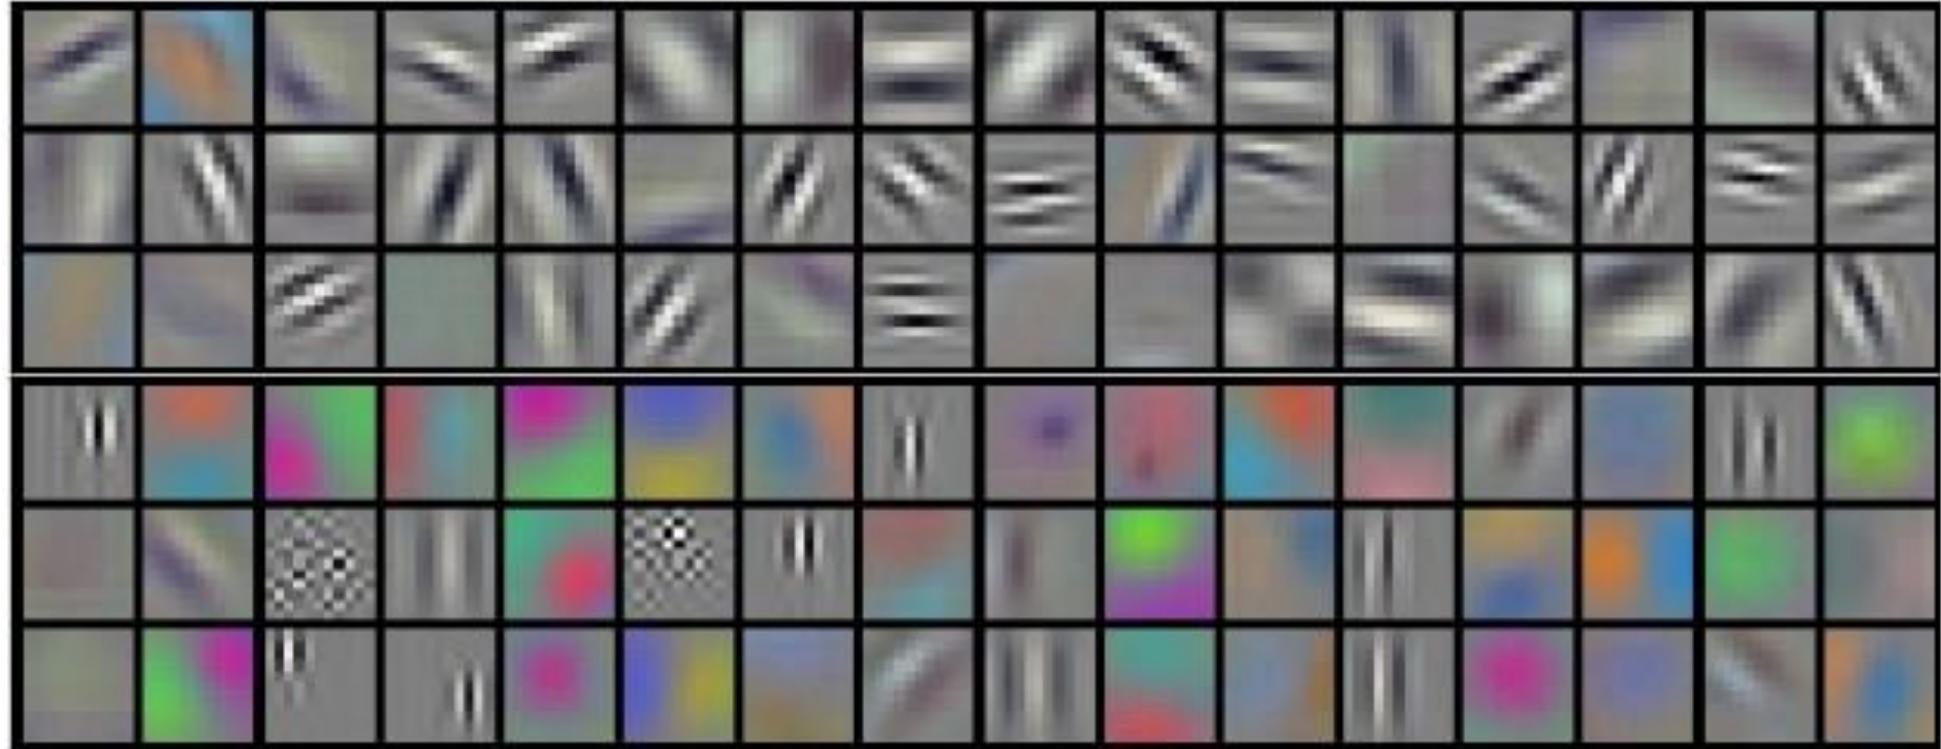


[Hands-On Machine Learning with Scikit-Learn and Tensorflow, Géron, 2017]

The Computational Graph Of A 1D Convolution

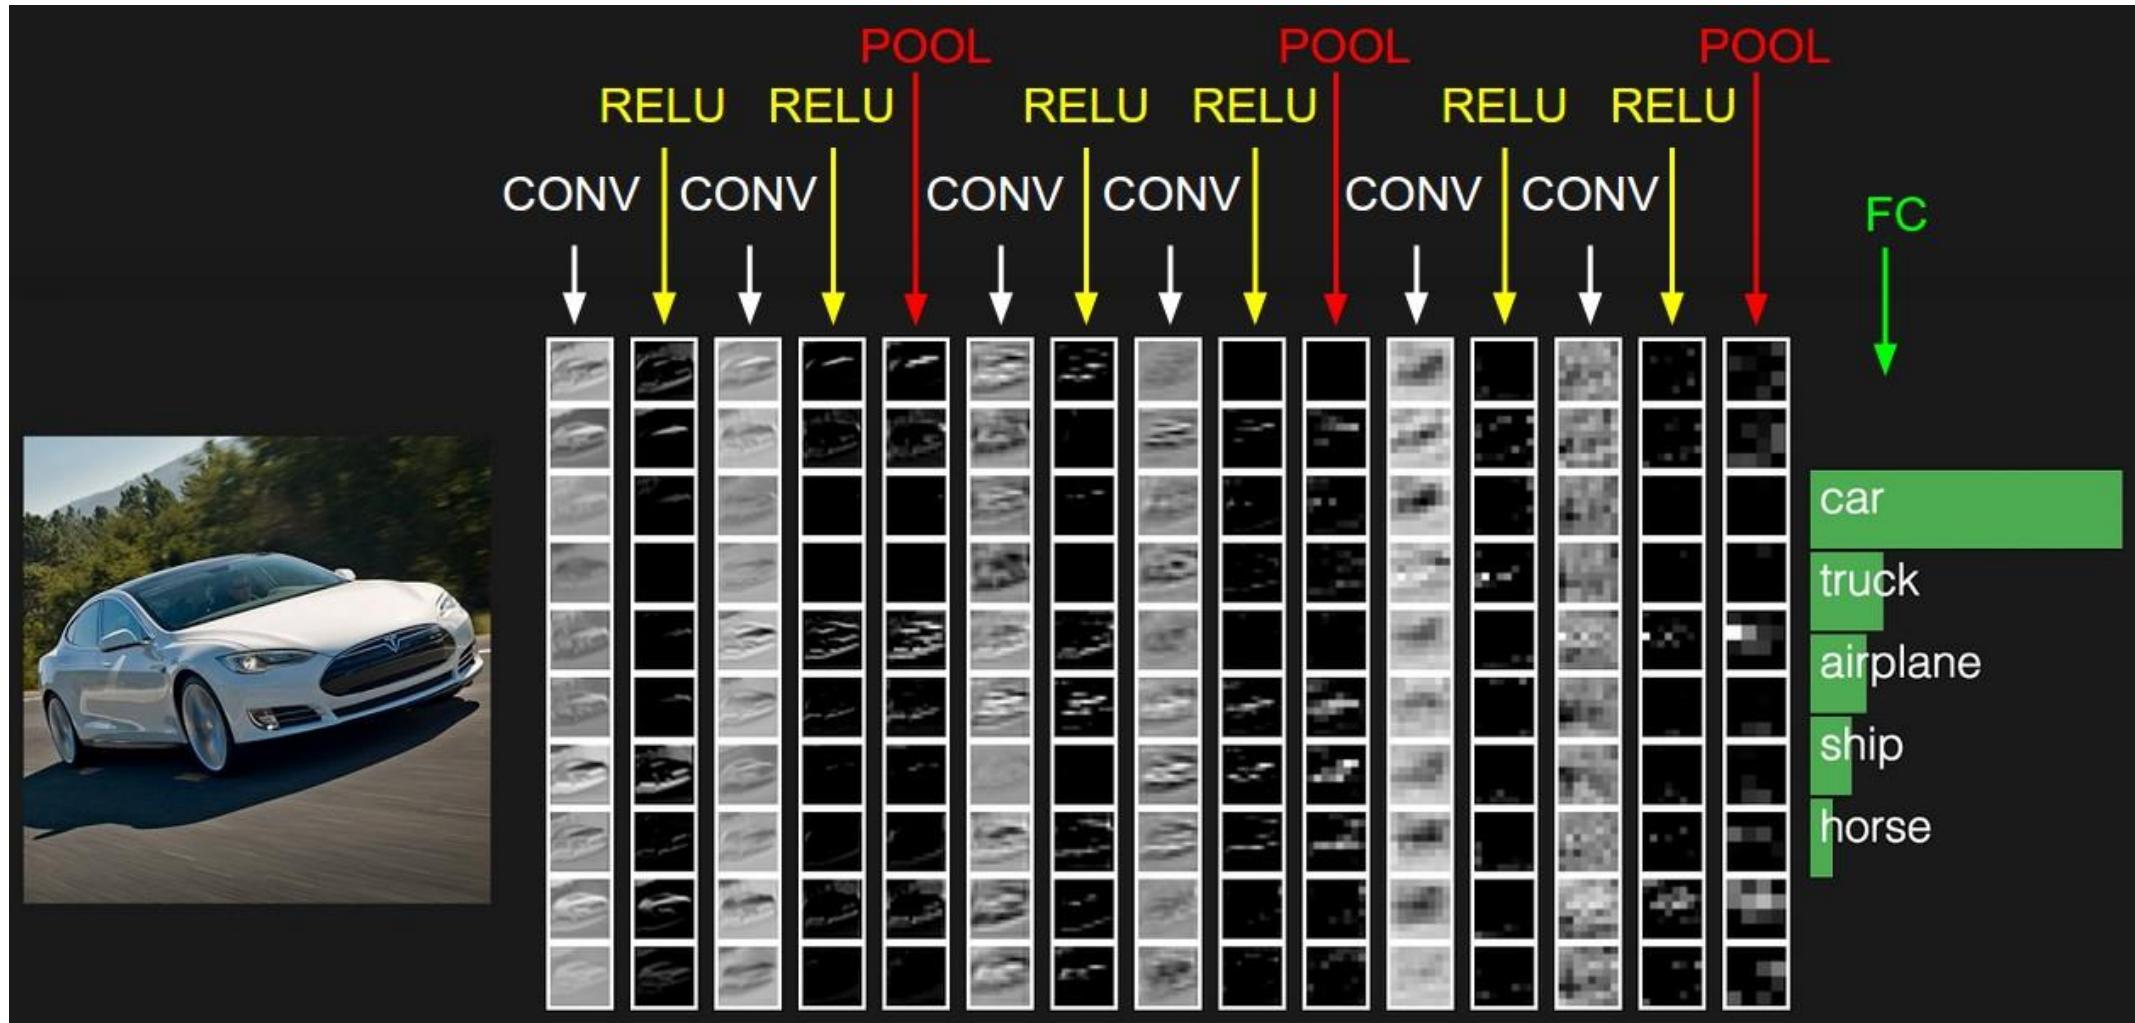


Learned Kernels



[Krizhevsky et al., 2012]

A CNN In Action



Implementation in PyTorch

```

import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

net = Net()
  
```

MaxPool2d ⚡

CLASS `torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False)`

[SOURCE]

Applies a 2D max pooling over an input signal composed of several input planes.

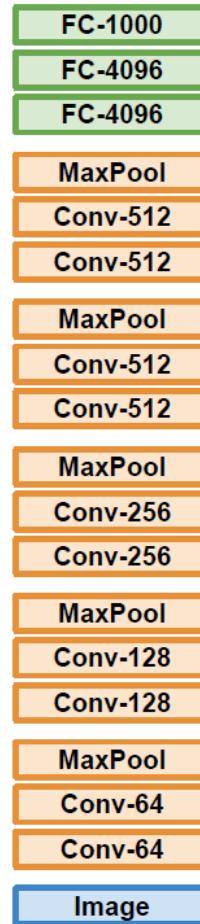
In the simplest case, the output value of the layer with input size (N, C, H, W) , output (N, C, H_{out}, W_{out}) and `kernel_size` (kH, kW) can be precisely described as:

$$out(N_i, C_j, h, w) = \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} \text{input}(N_i, C_j, \text{stride}[0] \times h + m, \text{stride}[1] \times w + n)$$

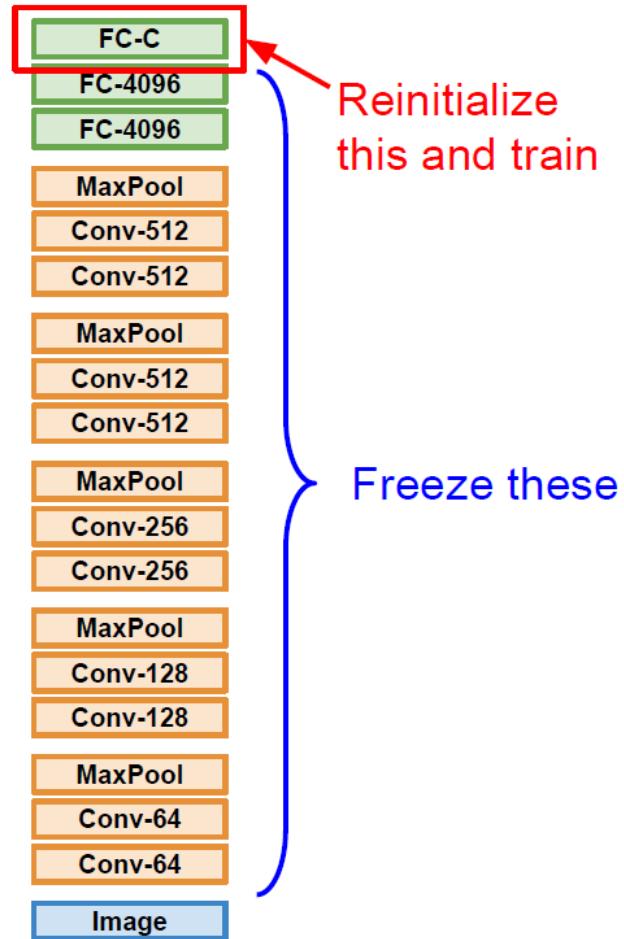
If `padding` is non-zero, then the input is implicitly zero-padded on both sides for `padding` number of points. `dilation` controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

Transfer Learning

1. Train on large data set
(or even pre-trained net)



2. Fine-tune: Train the last layer(s) for c custom classes



Later layers (e.g. the FC layers) act as good feature extractors (think very sophisticated basis function expansions).
Freezing means not updating the weights (even if you would get gradients from backprop).

```

import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)
alexnet = models.alexnet(pretrained=True)
squeezenet = models.squeezeNet1_0(pretrained=True)
vgg16 = models.vgg16(pretrained=True)
densenet = models.densenet161(pretrained=True)
inception = models.inception_v3(pretrained=True)
googlenet = models.googlenet(pretrained=True)
shufflenet = models.shufflenet_v2_x1_0(pretrained=True)
mobilenet = models.mobilenet_v2(pretrained=True)
resnext50_32x4d = models.resnext50_32x4d(pretrained=True)
wide_resnet50_2 = models.wide_resnet50_2(pretrained=True)
mnasnet = models.mnasnet1_0(pretrained=True)
  
```

Conclusion

Conclusion

Convolutional layers exploit sparse connectivity and parameter sharing to work with grid-structured data (images, time series)

Convolutional networks consist of at least a single convolutional layer

There is biological inspiration for that kind of local activation

Pooling layers perform subsampling